


Tutorial using the  software

---

## A tutorial for the R package `adeigenet_1.2-7`

T. JOMBART

---

### Looking for information?

More information is to be found from adegenet website: <http://adegenet.r-forge.r-project.org/>. Questions can be asked on the *adegenet forum* ([adegenet-forum@lists.r-forge.r-project.org](mailto:adegenet-forum@lists.r-forge.r-project.org)), a public mailing list whose archives are browsable and searchable. Please don't hesitate to use it! You will find more information about this forum in the section 'contact' of the adegenet website.

Comments and contributions on this tutorial are very welcome; please email me directly at: [t.jombart@imperial.ac.uk](mailto:t.jombart@imperial.ac.uk).

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>First steps</b>	<b>3</b>
2.1	Installing the package . . . . .	3
2.2	Object classes . . . . .	3
2.2.1	genind objects . . . . .	4
2.2.2	genpop objects . . . . .	8
<b>3</b>	<b>Various topics</b>	<b>9</b>
3.1	Importing data . . . . .	9
3.1.1	From GENETIX, STRUCTURE, FSTAT, Genepop . . . . .	9
3.1.2	From other software . . . . .	9
3.1.3	SNPs data . . . . .	11
3.1.4	DNA sequences . . . . .	11
3.1.5	Proteic sequences . . . . .	13
3.1.6	Using genind/genpop constructors . . . . .	17
3.2	Exporting data . . . . .	18
3.3	Manipulating data . . . . .	20
3.4	Using summaries . . . . .	25
3.5	Measuring and testing population structure (a.k.a F statistics) . . . . .	28
3.6	Testing for Hardy-Weinberg equilibrium . . . . .	31
3.7	Performing a Principal Component Analysis on <b>genind</b> objects . . . . .	32
3.8	Performing a Correspondance Analysis on <b>genpop</b> objects . . . . .	34
3.9	Analyzing a single locus . . . . .	36
3.10	Testing for isolation by distance . . . . .	38
3.11	Using Monmonier's algorithm to define genetic boundaries . . . . .	39
3.12	How to simulate hybridization? . . . . .	47
3.13	Handling presence/absence data . . . . .	49
3.14	Assigning genotypes to clusters using Discriminant Analysis . . . . .	52
3.14.1	Defining clusters . . . . .	52
3.14.2	Assigning new individuals . . . . .	59
<b>4</b>	<b>Frequently Asked Questions</b>	<b>62</b>
4.0.3	The function ... is not found. What's wrong? . . . . .	62

# 1 Introduction

This tutorial proposes a short visit through functionalities of the **adegenet** package for R (Ihaka & Gentleman, 1996; R Development Core Team, 2009). The purpose of this package is to facilitate the multivariate analysis of molecular marker data, especially using the **ade4** package (Chessel *et al.*, 2004). Data can be imported from a wide range of formats, including those of popular software (GENETIX, STRUCTURE, Fstat, Genepop), or from simple data frame of genotypes. **adegenet** also aims at providing a platform from which to use easily methods provided by other R packages (e.g., Goudet, 2005). Indeed, if it is possible to perform various genetic data analyses using R, data formats often differ from one package to another, and conversions are sometimes far from easy and straightforward.

In this tutorial, I first present the two object classes used in **adegenet**, namely **genind** (genotypes of individuals) and **genpop** (genotypes grouped by populations). Then, several topics will be tackled using reproducible examples.

## 2 First steps

### 2.1 Installing the package

Current version of the package is 1.2-3, and is compatible with R 2.8.1. Please make sure to be using at least R 2.8.1 and adegenet 1.2-3 before sending question about missing functions.

Here the **adegenet** package is installed along with other recommended packages.

```
> install.packages("adegenet", dep = TRUE)
```

Then the first step is to load the package:

```
> library(adegenet)
```

### 2.2 Object classes

Two classes of objects are defined, depending on the level at which the genetic information is stored: **genind** is used for individual genotypes, whereas **genpop** is used for alleles numbers counted by populations. Note that the term 'population', here and later, is employed in a broad sense: it simply refers to any grouping of individuals.

### 2.2.1 genind objects

These objects can be obtained by reading data files from other software, from a `data.frame` of genotypes, by conversion from a table of allelic frequencies, or even from aligned DNA sequences (see 'importing data').

```
> data(nancycats)
> is.genind(nancycats)
```

```
[1] TRUE
```

```
> nancycats
```

```
#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: genind(tab = truenames(nancycats)$tab, pop = truenames(nancycats)$pop)

@tab: 237 x 108 matrix of genotypes

@ind.names: vector of 237 individual names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 108 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: xy
```

A `genind` object is formal S4 object with several slots, accessed using the '@' operator (see `class?genind`). Note that the '\$' was also implemented for adegenet objects, so that slots can be accessed as if they were components of a list. The main slot in `genind` is a table of allelic frequencies of individuals (in rows) for every alleles in every loci. Being frequencies, data sum to one per locus, giving the score of 1 for an homozygote and 0.5 for an heterozygote. The particular case of presence/absence data will is described in an ad-hoc section (see 'Handling presence/absence data'). For instance:

```
> nancycats$tab[10:18, 1:10]
```

	L1.01	L1.02	L1.03	L1.04	L1.05	L1.06	L1.07	L1.08	L1.09	L1.10
010	0	0	0	0	0	0.0	0.0	0.0	1.0	0.0
011	0	0	0	0	0	0.0	0.0	0.0	0.0	0.5
012	0	0	0	0	0	0.5	0.0	0.5	0.0	0.0
013	0	0	0	0	0	0.5	0.0	0.5	0.0	0.0
014	0	0	0	0	0	0.0	0.0	1.0	0.0	0.0
015	0	0	0	0	0	0.0	0.5	0.0	0.5	0.0
016	0	0	0	0	0	0.5	0.0	0.0	0.5	0.0
017	0	0	0	0	0	0.5	0.0	0.5	0.0	0.0
018	0	0	0	0	0	0.5	0.0	0.0	0.5	0.0

Individual '010' is an homozygote for the allele 09 at locus 1, while '018' is an heterozygote with alleles 06 and 09. As user-defined labels are not always valid (for instance, they can be duplicated), generic labels are used for individuals, markers, alleles and eventually population. The true names are stored in the object (components `$[...].names` where ... can be 'ind', 'loc', 'all' or 'pop'). For instance :

```
> nancycats$loc.names
```

```
      L1      L2      L3      L4      L5      L6      L7      L8      L9
"fca8" "fca23" "fca43" "fca45" "fca77" "fca78" "fca90" "fca96" "fca37"
```

gives the true marker names, and

```
> nancycats$all.names[[3]]
```

```
      01      02      03      04      05      06      07      08      09      10
"133" "135" "137" "139" "141" "143" "145" "147" "149" "157"
```

gives the allele names for marker 3. Alternatively, one can use the accessor `locNames`:

```
> locNames(nancycats)
```

```
      L1      L2      L3      L4      L5      L6      L7      L8      L9
"fca8" "fca23" "fca43" "fca45" "fca77" "fca78" "fca90" "fca96" "fca37"
```

```
> head(locNames(nancycats, withAlleles = TRUE), 10)
```

```
[1] "fca8.117" "fca8.119" "fca8.121" "fca8.123" "fca8.127" "fca8.129"
[7] "fca8.131" "fca8.133" "fca8.135" "fca8.137"
```

The slot 'ploidy' is an integer giving the level of ploidy of the considered organisms (defaults to 2). This parameter is essential, in particular when switching from individual frequencies (genind object) to allele counts per populations (genpop). The slot 'type' describes the type of marker used: codominant ('codom', e.g. microsatellites) or presence/absence ('PA', e.g. AFLP). By default, adegenet considers that markers are codominant. Note that actual handling of presence/absence markers has been made available since version 1.2-3. See the dedicated section for more information about presence/absence markers.

Optional components are also allowed. The slot @other is a list that can include any additional information. The optional slot @pop (a factor giving a grouping of individuals) is particular in that the behaviour of many functions will check automatically for it and behave accordingly. In fact, each time an argument 'pop' is required by a function, it is first sought in @pop. For instance, using the function `genind2genpop` to convert `nancycats` to a `genpop` object, there is no need to give a 'pop' argument as it exists in the `genind` object:

```
> table(nancycats$pop)
```

```
P01 P02 P03 P04 P05 P06 P07 P08 P09 P10 P11 P12 P13 P14 P15 P16 P17
 10  22  12  23  15  11  14  10   9  11  20  14  13  17  11  12  13
```

```
> catpop <- genind2genpop(nancycats)
```

```
Converting data from a genind to a genpop object...
```

```
...done.
```

```
> catpop
```

```
#####
### Genpop object ###
#####
- Alleles counts for populations -

S4 class:   genpop
@call:    genind2genpop(x = nancycats)

@tab:    17 x 108 matrix of alleles counts

@pop.names: vector of 17 population names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 108 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type: codom

@other: a list containing: xy
```

Other additional components can be stored (like here, spatial coordinates of populations in `$xy`) but will not be passed during any conversion (`catpop` has no `$other$xy`).

Note that the slot 'pop' can be retrieved and set using the `pop` function:

```
> obj <- nancycats[sample(1:50, 10)]
> pop(obj)
```

```
[1] 3 2 2 4 4 3 1 2 2 3
Levels: 3 2 4 1
```

```
> pop(obj) <- rep("newPop", 10)
> pop(obj)
```

```
[1] newPop newPop newPop newPop newPop newPop newPop newPop newPop newPop
Levels: newPop
```

Finally, a `genind` object generally contains its matched call, *i.e.* the instruction that created it. This is not the case, however, for objects loaded using `data`. When call is available, it can be used to regenerate an object.

```
> obj <- read.genetix(system.file("files/nancycats.gtx", package = "adegenet"))
```

```
Converting data from GENETIX to a genind object...
...done.
```

```
> obj$call
```

```
read.genetix(file = system.file("files/nancycats.gtx", package = "adegenet"))
```

```
> toto <- eval(obj$call)
```

```
Converting data from GENETIX to a genind object...
...done.
```

```
> identical(obj, toto)
```

```
[1] TRUE
```

## 2.2.2 genpop objects

We use the previously built `genpop` object:

```
> catpop
```

```
#####
### Genpop object ###
#####
- Alleles counts for populations -

S4 class:   genpop
@call:   genind2genpop(x = nancycats)

@tab:   17 x 108 matrix of alleles counts

@pop.names: vector of 17 population names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 108 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type: codom

@other: a list containing: xy
```

```
> is.genpop(catpop)
```

```
[1] TRUE
```

```
> catpop$tab[1:5, 1:10]
```

	L1.01	L1.02	L1.03	L1.04	L1.05	L1.06	L1.07	L1.08	L1.09	L1.10
01	0	0	0	0	0	0	0	2	9	1
02	0	0	0	0	0	10	9	8	14	2
03	0	0	0	4	0	0	0	0	1	10
04	0	0	0	3	0	0	0	1	7	17
05	0	0	0	1	0	0	0	0	7	10

The matrix `$tab` contains alleles counts per population (here, cat colonies). These objects are otherwise very similar to `genind` in their structure, and possess generic names, true names, the matched call and an `@other` slot.



## 3 Various topics

### 3.1 Importing data

#### 3.1.1 From GENETIX, STRUCTURE, FSTAT, Genepop

Data can be read from the software GENETIX (.gtx), STRUCTURE (.str or .stru), FSTAT (.dat) and Genepop (.gen) files, using the corresponding `read` function: `read.genetix`, `read.structure`, `read.fstat`, and `read.genepop`. These functions take as main argument the path (as a string character) to an input file, and produce a `genind` object. Alternatively, one can use the function `import2genind` which detects a file format from its extension and uses the appropriate routine. For instance:

```
> obj1 <- read.genetix(system.file("files/nancycats.gtx", package = "adegenet"))

Converting data from GENETIX to a genind object...
...done.

> obj2 <- import2genind(system.file("files/nancycats.gtx", package = "adegenet"))

Converting data from GENETIX to a genind object...
...done.

> all.equal(obj1, obj2)

[1] "Attributes: < Component 2: target, current do not match when deparsed >"
```

The only difference between `obj1` and `obj2` is their call (which is normal as they were obtained from different command lines).

#### 3.1.2 From other software

Genetic markers data can most of the time be stored as a table with individuals in row and markers in column, where each entry is a character string coding the alleles possessed at one locus. Such data are easily imported into R as a `data.frame`, using for instance `read.table` for text files or `read.csv` for comma-separated text files. Then, the obtained `data.frame` can be converted into a `genind` object using `df2genind`.

There are only a few pre-requisite the data should meet for this conversion to be possible. The easiest and clearest way of coding data is using a separator

between alleles. For instance, "80/78", "80|78", or "80,78" are different ways of coding a genotype at a microsatellite locus with alleles '80' and '78'. Note that for haploid data, no separator shall be used. As a consequence, SNP data should consist of the raw nucleotides. The only constraint when using a separator is that the same separator is used in all the dataset. There are no constraints as to i) the type of separator used or ii) the ploidy of the data. These parameters can be set in **df2genind** through arguments 'sep' and 'ploidy', respectively.

Alternatively, no separator may be used provided a fixed number of characters is used to code any allele. For instance, in a diploid organism, "0101" is an homozygote 1/1 while "1209" is a heterozygote 12/09 in a two-character per allele coding scheme. In a tetraploid system with one character per allele, "1209" will be understood as 1/2/0/9.

Here, I provide an example using a data set from the library hierfstat.

```
> library(hierfstat)
> toto <- read.fstat.data(paste(.path.package("hierfstat"), "/data/diploid.dat",
+   sep = "", collapse = ""), nloc = 5)
> head(toto)
```

	Pop	loc-1	loc-2	loc-3	loc-4	loc-5
1	1	44	43	43	33	44
2	1	44	44	43	33	44
3	1	44	44	43	43	44
4	1	44	44	NA	33	44
5	1	44	44	24	34	44
6	1	44	44	NA	43	44

**toto** is a data frame containing genotypes and a population factor.

```
> obj <- df2genind(X = toto[, -1], pop = toto[, 1])
> obj
```

```
#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:   genind
@call: df2genind(X = toto[, -1], pop = toto[, 1])

@tab:  44 x 11 matrix of genotypes

@ind.names: vector of  44 individual names
@loc.names: vector of  5 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the  11 columns of @tab
@all.names: list of  5 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: - empty -
```

`obj` is a **genind** containing the same information, but recoded as a matrix of allele frequencies (`$tab` slot).

### 3.1.3 SNPs data

In *adegenet*, SNP data are handled as other codominant markers such as microsatellites. The most convenient way to convert SNPs into a **genind** is using **df2genind**, which is described in the previous section. Let `dat` be an input matrix, as can be read into R using **read.table** or **read.csv**, with genotypes in row and SNP loci in columns.

```
> dat <- matrix(sample(c("a", "t", "g", "c"), 15, replace = TRUE),
+               nrow = 3)
> rownames(dat) <- paste("genot.", 1:3)
> colnames(dat) <- 1:5
> dat
```

```
      1 2 3 4 5
genot. 1 "g" "t" "c" "g" "t"
genot. 2 "g" "a" "t" "c" "a"
genot. 3 "t" "g" "c" "a" "g"
```

```
> obj <- df2genind(dat, ploidy = 1)
> truenames(obj)
```

```
      1.g 1.t 2.a 2.g 2.t 3.c 3.t 4.a 4.c 4.g 5.a 5.g 5.t
genot. 1 1 0 0 0 1 1 0 0 1 0 0 1
genot. 2 1 0 1 0 0 0 1 1 0 1 0 0
genot. 3 0 1 0 1 0 1 0 1 0 0 1 0
```

`obj` is a **genind** containing the SNPs information, which can be used for further analysis in *adegenet*.

### 3.1.4 DNA sequences

DNA sequences can be read into R using the *ape* package (Paradis *et al.*, 2004; Paradis, 2006), and imported into *adegenet* using **DNAbin2genind**. There are several ways *ape* can be used to read in DNA sequences. The easiest one is reading data from a usual format such as FASTA or Clustal using **read.dna**. Other options include reading data directly from GenBank using **read.GenBank**, or from other public databases using the *seqinr* package and transforming the **alignment** object into a **DNAbin** using **as.DNAbin**.

Here, we illustrate this approach by re-using the example of **read.GenBank**. A connection to the internet is required, as sequences are read directly from a distant database.

```
> library(ape)
> ref <- c("U15717", "U15718", "U15719", "U15720", "U15721", "U15722",
+         "U15723", "U15724")
> myDNA <- read.GenBank(ref)
> myDNA
```

8 DNA sequences in binary format stored in a list.

All sequences of same length: 1045

Labels: U15717 U15718 U15719 U15720 U15721 U15722 ...

Base composition:

a	c	g	t
0.267	0.351	0.134	0.247

```
> class(myDNA)
```

```
[1] "DNAbin"
```

```
> summary(myDNA)
```

	Length	Class	Mode
U15717	1045	DNAbin	raw
U15718	1045	DNAbin	raw
U15719	1045	DNAbin	raw
U15720	1045	DNAbin	raw
U15721	1045	DNAbin	raw
U15722	1045	DNAbin	raw
U15723	1045	DNAbin	raw
U15724	1045	DNAbin	raw

In *adegenet*, only polymorphic loci are conserved; importing data from a DNA sequence to *adegenet* therefore consist in extracting SNPs from the aligned sequences. This conversion is achieved by **DNAbin2genind**. This function allows one to specify a threshold for polymorphism; for instance, one could retain only SNPs for which the second largest allele frequency is greater than 1% (using the **polyThres** argument). This is achieved using:

```
> obj <- DNAbin2genind(myDNA, polyThres = 0.01)
> obj

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: DNAbin2genind(x = myDNA, polyThres = 0.01)

@tab: 8 x 318 matrix of genotypes
```

```
@ind.names: vector of 8 individual names
@loc.names: vector of 155 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 318 columns of @tab
@all.names: list of 155 components yielding allele names for each locus
@ploidy: 1
@type: codon

Optionnal contents:
@pop: - empty -
@pop.names: - empty -

@other: - empty -
```

Here, out of the 1045 nucleotides of the sequences, 318 SNPs were extracted and stored as a **genind** object.

### 3.1.5 Proteic sequences

Alignments of proteic sequences can be exploited in *adegenet* in the same way as DNA sequences (see section above). Alignments are scanned for polymorphic sites, and only those are retained to form a **genind** object. Loci correspond to the position of the residue in the alignment, and alleles correspond to the different amino-acids (AA). Aligned proteic sequences are stored as objects of class **alignment** in the *seqinr* package (Charif & Lobry, 2007). See **?as.alignment** for a description of this class. The function extracting polymorphic sites from **alignment** objects is **alignment2genind**.

Its use is fairly simple. It is here illustrated using a small dataset of aligned proteic sequences:

```
> library(seqinr)
> mase.res <- read.alignment(file = system.file("sequences/test.mase",
+ package = "seqinr"), format = "mase")
> mase.res

$nb
[1] 6

$nam
[1] "Langur" "Baboon" "Human" "Rat" "Cow" "Horse"

$seq
$seq[[1]]
[1] "-kifercelartlkklgldgykgvslanwvclakwesgynteatnynpgdestdygifqinsrywcnnkpgavdachiscsallqnniada"

$seq[[2]]
[1] "-kifercelartlkrllgldgyrgislanwvclakwesdyntqatnynpgdqstdygifqinshywcndgkpgavnachiscnallqdnitda"

$seq[[3]]
[1] "-kvfercelartlkrllgmdgyrgislanwvclakwesgyntnatnynagdrstdygifqinsrywcndgkpgavnachlscsallqdniada"

$seq[[4]]
[1] "-ktyercefartlkrngmsgygysladwvclaqhesnyntqarnydpdqstdygifqinsrywcndgkpraknacgipcsallqdditqa"
```

```
$seq[[5]]
[1] "-kvfercelartlkklgldgkgyvlanwlccltkwessyntkatnynpssestdygifqinskwncndgkpnavgchvscselmendiaka

$seq[[6]]
[1] "-kvfiskcelahklkaqemdggyslanwvcmaeyesnfnttrafngknangssdyglfqlnnkwckdnkrsssnacnimcsklldenidd

$com
[1] ";empty description\n" ";\\n"
[4] ";\\n" ";\\n"

attr("class")
[1] "alignment"
```

```
> x <- alignment2genind(mase.res)
> x
```

```
#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: alignment2genind(x = mase.res)

@tab: 6 x 212 matrix of genotypes

@ind.names: vector of 6 individual names
@loc.names: vector of 82 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 212 columns of @tab
@all.names: list of 82 components yielding allele names for each locus
@ploidy: 1
@type: codon

Optionnal contents:
@pop: - empty -
@pop.names: - empty -

@other: a list containing: com
```

The six aligned protein sequences (`mase.res`) have been scanned for polymorphic sites, and these have been extracted to form the `genind` object `x`. Note that several settings such as the characters corresponding to missing values (i.e., gaps) and the for polymorphism threshold for a site to be retained can be specified through the function's arguments (see `?alignment2genind`).

The names of the loci directly provides the indices of polymorphic sites:

```
> locNames(x)
```

L01	L02	L03	L04	L05	L06	L07	L08	L09	L10	L11	L12	L13
"3"	"4"	"5"	"6"	"9"	"11"	"12"	"15"	"16"	"17"	"18"	"19"	"21"
L14	L15	L16	L17	L18	L19	L20	L21	L22	L23	L24	L25	L26
"22"	"24"	"28"	"30"	"32"	"33"	"34"	"35"	"38"	"39"	"42"	"44"	"46"
L27	L28	L29	L30	L31	L32	L33	L34	L35	L36	L37	L38	L39

```
"47"  "48"  "49"  "50"  "51"  "53"  "57"  "60"  "62"  "63"  "64"  "67"  "68"
L40   L41   L42   L43   L44   L45   L46   L47   L48   L49   L50   L51   L52
"69"  "71"  "72"  "73"  "74"  "75"  "76"  "78"  "79"  "80"  "82"  "83"  "85"
L53   L54   L55   L56   L57   L58   L59   L60   L61   L62   L63   L64   L65
"86"  "87"  "88"  "90"  "91"  "92"  "93"  "94"  "98"  "99"  "101" "102" "103"
L66   L67   L68   L69   L70   L71   L72   L73   L74   L75   L76   L77   L78
"105" "106" "109" "112" "113" "114" "116" "117" "118" "120" "121" "122" "124"
L79   L80   L81   L82
"125" "126" "128" "129"
```

The table of polymorphic sites can be reconstructed easily by:

```
> tabAA <- genind2df(x)
> dim(tabAA)
```

```
[1] 6 82
```

```
> tabAA[, 1:20]
```

```
      3 4 5 6 9 11 12 15 16 17 18 19 21 22 24 28 30 32 33 34
Langur i f e r l r t k l g l d y k v n v l a k
Baboon i f e r l r t r l g l d y r i n v l a k
Human  v f e r l r t r l g m d y r i n m l a k
Rat    t y e r f r t r n g m s y y v d v l a q
Cow    v f e r l r t k l g l d y k v n l l t k
Horse  v f s k l h k a q e m d f g y n v m a e
```

The global AA composition of the polymorphic sites is given by:

```
> table(unlist(tabAA))
```

```
 a d e f g h i k l m n p q r s t v w y
35 38 16 9 33 13 27 28 31 8 44 10 26 47 36 20 42 6 23
```


Now that polymorphic sites have been converted into a `genind` object, simple distances can be computed between the sequences. Note that `adeigenet` does not implement specific distances for protein sequences, we only use the simple Euclidean distance. Fancier protein distances are implemented in R; see for instance `dist.alignment` in the *seqinr* package, and `dist.ml` in the *phangorn* package.

```
> D <- dist(truenames(x))
> D
```

```
      Langur  Baboon  Human  Rat  Cow
Baboon 5.291503
Human  6.000000 5.291503
Rat    8.717798 8.124038 8.602325
Cow    7.874008 8.717798 8.944272 10.392305
Horse  11.313708 11.313708 11.224972 11.224972 11.747340
```

This matrix of distances is small enough for one to interpret the raw numbers. However, it is also very straightforward to represent these distances as a tree or in a reduced space. We first build a Neighbor-Joining tree using the *ape* package:

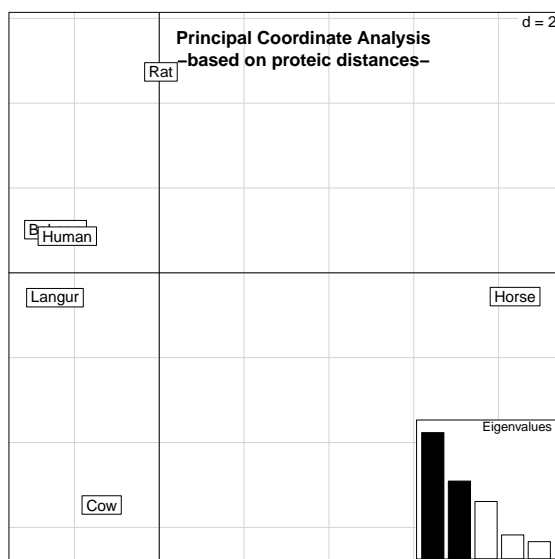
```
> library(ape)
> tre <- nj(D)
> plot(tre, type = "unrooted", edge.w = 2)
> edgelabels(tex = round(tre$edge.length, 1), bg = rgb(0.8, 0.8,
+ 1, 0.8))
```



The best possible planar representation of these Euclidean distances is achieved by Principal Coordinate Analyses (PCoA), which in this case will give identical results to PCA of the original (centred, non-scaled) data:

```
> pco1 <- dudi.pco(D, scannf = FALSE, nf = 2)
> scatter(pco1, posi = "bottomright")
> title("Principal Coordinate Analysis\nn-based on proteic distances-")
```





### 3.1.6 Using `genind`/`genpop` constructors

Lastly, `genind` or `genpop` objects can be constructed from data matrices similar to the `$tab` component (respectively, alleles frequencies and alleles counts). This is achieved by the constructors `genind` (or `as.genind`) and `genpop` (or `as.genpop`). However, these low-level functions are first meant for internal use, and are called for instance by functions such as `read.genetix`. Consequently, there is much less control on the arguments and improper specification can lead to creating improper `genind`/`genpop` objects without issuing a warning or an error, by leading to meaningless subsequent analysis.

Therefore, one should use these functions with additional care as to how information is coded. The table passed as argument to these constructors must have correct names: unique rownames identifying genotypes/populations, and unique colnames having the form `'[marker].[allele]'`.

Here is an example for `genpop` using a dataset from `ade4`:

```
> library(ade4)
> data(microsatt)
> microsatt$tab[10:15, 12:15]
```

	INRA32.168	INRA32.170	INRA32.174	INRA32.176
Mtbeliard	0	0	0	1
NDama	0	0	0	12
Normand	1	0	0	2
Parthenais	8	5	0	3
Somba	0	0	0	20
Vosgienne	2	0	0	0

`microsat$tab` contains alleles counts per populations, and can therefore be used to make a `genpop` object. Moreover, column names are set as required, and row names are unique. It is therefore safe to convert these data into a `genpop` using the constructor:

```
> toto <- genpop(microsat$tab)
> toto

#####
### Genpop object ###
#####
- Alleles counts for populations -

S4 class:   genpop
@call:   genpop(tab = microsat$tab)

@tab:   18 x 112 matrix of alleles counts

@pop.names: vector of 18 population names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 112 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type: codom

@other: - empty -

> summary(toto)

# Number of populations: 18

# Number of alleles per locus:
L1 L2 L3 L4 L5 L6 L7 L8 L9
8 15 11 10 17 10 14 15 12

# Number of alleles per population:
01 02 03 04 05 06 07 08 09 10 11 12 13 14 15 16 17 18
39 69 51 59 52 41 34 48 46 47 43 56 57 52 49 64 56 67

# Percentage of missing data:
[1] 0
```

## 3.2 Exporting data

Genotypes in `genind` format can be exported to the R packages *genetics* (using `genind2genotype`) and *hierfstat* (using `genind2hierfstat`). The package *genetics* is now deprecated, but the implemented class `genotype` is still used in various packages. The package *hierfstat* does not define a class, but requires data to be formatted in a particular way. Here are examples of how to use these functions:

```
> obj <- genind2genotype(nancycats)
> class(obj)

[1] "data.frame"

> obj[1:4, 1:5]

      fca8  fca23  fca43  fca45  fca77
N215  <NA> 136/146 139/139 120/116 156/156
N216  <NA> 146/146 139/145 126/120 156/156
N217 135/143 136/146 141/141 116/116 156/152
N218 135/133 138/138 139/141 126/116 150/150
```

```
> class(obj$fca8)

[1] "genotype" "factor"
```

```
> obj <- genind2hierfstat(nancycats)
> class(obj)

[1] "data.frame"
```

```
> obj[1:4, 1:5]

      pop  fca8  fca23  fca43  fca45
N215   1    NA 136146 139139 116120
N216   1    NA 146146 139145 120126
N217   1 135143 136146 141141 116116
N218   1 133135 138138 139141 116126
```

Now we can use the function `varcomp.glob` from *hierfstat* to compute 'variance' components:

```
> varcomp.glob(obj$pop, obj[, -1])

$loc
      [,1]      [,2]      [,3]
fca8 0.08867161 0.116693199 0.6682028
fca23 0.05384247 0.077539920 0.6666667
fca43 0.05518935 0.066055996 0.6793249
fca45 0.05861271 -0.001026783 0.7083333
fca77 0.08810966 0.156863586 0.6329114
fca78 0.04869695 0.079006911 0.5654008
fca90 0.07540329 0.097194716 0.6497890
fca96 0.07538325 -0.005902071 0.7543860
fca37 0.04264094 0.116318729 0.4514768

$overall
      Pop      Ind      Error
0.5865502 0.7027442 5.7764917

$F
      Pop      Ind
Total 0.08301274 0.1824701
Pop    0.00000000 0.1084610
```

A more generic way to export data is to produce a data.frame of genotypes coded by character strings. This is done by `genind2df`:

```
> obj <- genind2df(nancycats)
> obj[1:5, 1:5]
```

	pop	fca8	fca23	fca43	fca45
N215	1	<NA>	136146	139139	116120
N216	1	<NA>	146146	139145	120126
N217	1	135143	136146	141141	116116
N218	1	133135	138138	139141	116126
N219	1	133135	140146	141145	126126

However, some software will require alleles to be separated. The argument `sep` allows one to specify any separator. For instance:

```
> genind2df(nancycats, sep = "|")[1:5, 1:5]
```

	pop	fca8	fca23	fca43	fca45
N215	1	<NA>	136 146	139 139	116 120
N216	1	<NA>	146 146	139 145	120 126
N217	1	135 143	136 146	141 141	116 116
N218	1	133 135	138 138	139 141	116 126
N219	1	133 135	140 146	141 145	126 126

Note that tabulations can be obtained as follows using `'\t'` character.

### 3.3 Manipulating data

Data manipulation is meant to be easy in `adegenet` (if it is not, complain!). First, as `genind` and `genpop` objects are basically formed by a data matrix (the `@tab` slot), it is natural to subset these objects like it is done with a matrix. The `[]` operator does this, forming a new object with the retained genotypes/populations and alleles:

```
> titi <- toto[1:3, ]
> toto$pop.names
```

	01	02	03	04	05	06
"Baoule"		"Borgou"	"BPN"	"Charolais"	"Holstein"	"Jersey"
	07	08	09	10	11	12
"Lagunaire"		"Limousin"	"MaineAnjou"	"Mtbeliard"	"NDama"	"Normand"
	13	14	15	16	17	18
"Parthenais"		"Somba"	"Vosgienne"	"ZChoa"	"ZMbororo"	"Zpeul"

```
> titi
```

```
#####  
### Genpop object ###  
#####  
- Alleles counts for populations -  
  
S4 class:  genpop  
@call:  .local(x = x, i = i, j = j, drop = drop)  
  
@tab:  3 x 112 matrix of alleles counts  
  
@pop.names: vector of  3 population names  
@loc.names: vector of  9 locus names  
@loc.nall: number of alleles per locus  
@loc.fac: locus factor for the  112 columns of @tab  
@all.names: list of  9 components yielding allele names for each locus  
@ploidy:  2  
@type:  codom  
  
@other: a list containing: elements without names  
  
> titi$pop.names  
  
      1      2      3  
"Baoule" "Borgou" "BPN"
```

The object **toto** has been subsetted, keeping only the first three populations. Of course, any subsetting available for a matrix can be used with **genind** and **genpop** objects. For instance, we can subset **titi** to keep only the third marker:

```
> titi <- titi[, titi$loc.fac == "L3"]  
> titi  
  
#####  
### Genpop object ###  
#####  
- Alleles counts for populations -  
  
S4 class:  genpop  
@call:  .local(x = x, i = i, j = j, drop = drop)  
  
@tab:  3 x 11 matrix of alleles counts  
  
@pop.names: vector of  3 population names  
@loc.names: vector of  1 locus names  
@loc.nall: number of alleles per locus  
@loc.fac: locus factor for the  11 columns of @tab  
@all.names: list of  1 components yielding allele names for each locus  
@ploidy:  2  
@type:  codom  
  
@other: a list containing: elements without names
```

Now, **titi** only contains the 11 alleles of the third marker of **toto**.

To simplify the task of separating data by marker, the function **seploc** can be used. It returns a list of objects (optionnaly, of data matrices), each corresponding to a marker:

```
> sepCats <- seploc(nancycats)
> class(sepCats)

[1] "list"

> names(sepCats)

[1] "fca8" "fca23" "fca43" "fca45" "fca77" "fca78" "fca90" "fca96" "fca37"

> sepCats$fca45

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: .local(x = x)

@tab: 237 x 9 matrix of genotypes

@ind.names: vector of 237 individual names
@loc.names: vector of 1 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 9 columns of @tab
@all.names: list of 1 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: xy
```

The object `sepCats$fca45` only contains data of the marker fca45.

Following the same idea, `seppop` allows one to separate genotypes in a `genind` object by population. For instance, we can separate genotype of cattles in the dataset `microbov` by breed:

```
> data(microbov)
> obj <- seppop(microbov)
> class(obj)

[1] "list"

> names(obj)
```

```
[1] "Borgou"      "Zebu"        "Lagunaire"   "NDama"
[5] "Somba"      "Aubrac"      "Bazadais"    "BlondeAquitaine"
[9] "BretPieNoire" "Charolais"   "Gascon"      "Limousin"
[13] "MaineAnjou"  "Montbeliard" "Salers"
```

```
> obj$Borgou
```

```
#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:  genind
@call:  .local(x = x, i = i, j = j, treatOther = ..1, drop = drop)

@tab:  50 x 373 matrix of genotypes

@ind.names: vector of  50 individual names
@loc.names: vector of  30 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the  373 columns of @tab
@all.names: list of  30 components yielding allele names for each locus
@ploidy:  2
@type:  codom

Optionnal contents:
@pop:  factor giving the population of each individual
@pop.names:  factor giving the population of each individual

@other: a list containing: coun  breed  spe
```

The returned object `obj` is a list of `genind` objects each containing genotypes of a given breed.

A last, rather vicious trick is to separate data by population and by marker. This is easy using `lapply`; one can first separate population then markers, or the contrary. Here, we separate markers inside each breed in `obj`

```
> obj <- lapply(obj, seploc)
> names(obj)
```

```
[1] "Borgou"      "Zebu"        "Lagunaire"   "NDama"
[5] "Somba"      "Aubrac"      "Bazadais"    "BlondeAquitaine"
[9] "BretPieNoire" "Charolais"   "Gascon"      "Limousin"
[13] "MaineAnjou"  "Montbeliard" "Salers"
```

```
> class(obj$Borgou)
```

```
[1] "list"
```

```
> names(obj$Borgou)
```

```
[1] "INRA63" "INRA5" "ETH225" "ILSTS5" "HEL5" "HEL1" "INRA35"
[8] "ETH152" "INRA23" "ETH10" "HEL9" "CSSM66" "INRA32" "ETH3"
[15] "BM2113" "BM1824" "HEL13" "INRA37" "BM1818" "ILSTS6" "MM12"
[22] "CSRM60" "ETH185" "HAUT24" "HAUT27" "TGLA227" "TGLA126" "TGLA122"
[29] "TGLA53" "SPS115"
```

```
> obj$Borgou$INRA63
```

```
#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: .local(x = x)

@tab: 50 x 9 matrix of genotypes

@ind.names: vector of 50 individual names
@loc.names: vector of 1 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 9 columns of @tab
@all.names: list of 1 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: coun breed spe
```

For instance, `obj$Borgou$INRA63` contains genotypes of the breed Borgou for the marker INRA63.

Lastly, one may want to pool genotypes in different datasets, but having the same markers, into a single dataset. This is more than just merging the `@tab` components of all datasets, because alleles can differ (they almost always do) and markers are not necessarily sorted the same way. The function `repool` is designed to avoid these problems. It can merge any `genind` provided as arguments as soon as the same markers are used. For instance, it can be used after a `seppop` to retain only some populations:

```
> obj <- seppop(microbov)
> names(obj)
```

```
[1] "Borgou" "Zebu" "Lagunaire" "NDama"
[5] "Somba" "Aubrac" "Bazadais" "BlondeAquitaine"
[9] "BretPieNoire" "Charolais" "Gascon" "Limousin"
[13] "MaineAnjou" "Montbeliard" "Salers"
```



```
> newObj <- repool(obj$Borgou, obj$Charolais)
> newObj

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class: genind
@call: repool(obj$Borgou, obj$Charolais)

@tab: 105 x 295 matrix of genotypes

@ind.names: vector of 105 individual names
@loc.names: vector of 30 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 295 columns of @tab
@all.names: list of 30 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: - empty -

> newObj$pop.names

      P1      P2
"Borgou" "Charolais"
```

Done !

### 3.4 Using summaries

Both **genind** and **genpop** objects have a summary providing basic information about data. Informations are both printed and invisibly returned as a list.

```
> toto <- summary(nancycats)

# Total number of genotypes: 237

# Population sample sizes:
 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17
10 22 12 23 15 11 14 10  9 11 20 14 13 17 11 12 13

# Number of alleles per locus:
L1 L2 L3 L4 L5 L6 L7 L8 L9
16 11 10  9 12  8 12 12 18

# Number of alleles per population:
01 02 03 04 05 06 07 08 09 10 11 12 13 14 15 16 17
36 53 50 67 48 56 42 54 43 46 70 52 44 61 42 40 35
```

```
# Percentage of missing data:
[1] 2.344116

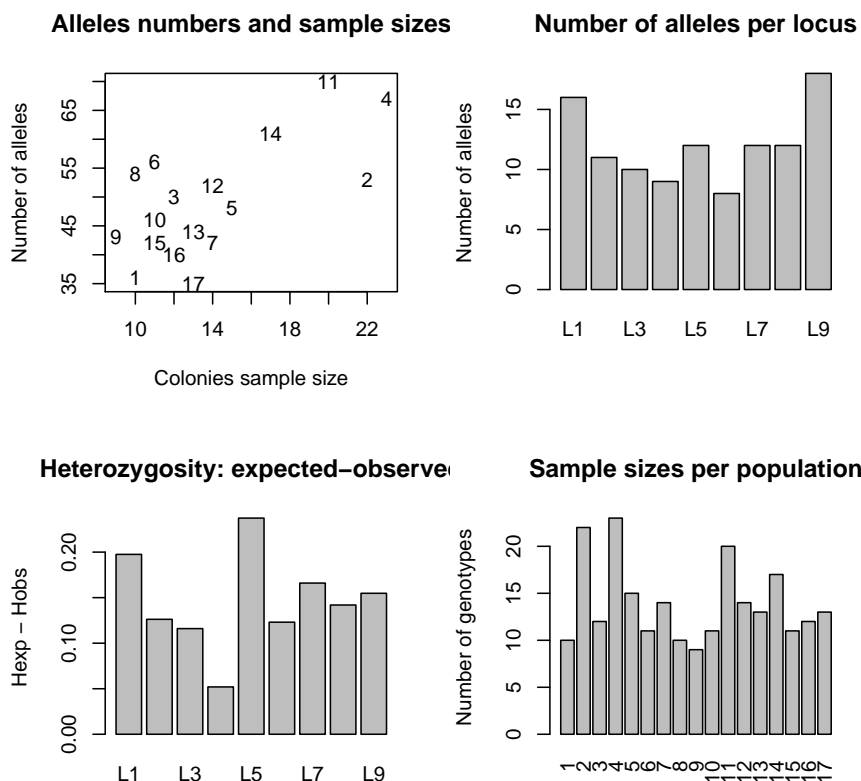
# Observed heterozygosity:
      L1      L2      L3      L4      L5      L6      L7      L8
0.6682028 0.6666667 0.6793249 0.7083333 0.6329114 0.5654008 0.6497890 0.6184211
      L9
0.4514768

# Expected heterozygosity:
      L1      L2      L3      L4      L5      L6      L7      L8
0.8657224 0.7928751 0.7953319 0.7603095 0.8702576 0.6884669 0.8157881 0.7603493
      L9
0.6062686

> names(toto)

[1] "N"      "pop.eff" "loc.nall" "pop.nall" "NA.perc" "Hobs"      "Hexp"

> par(mfrow = c(2, 2))
> plot(toto$pop.eff, toto$pop.nall, xlab = "Colonies sample size",
+      ylab = "Number of alleles", main = "Alleles numbers and sample sizes",
+      type = "n")
> text(toto$pop.eff, toto$pop.nall, lab = names(toto$pop.eff))
> barplot(toto$loc.nall, ylab = "Number of alleles", main = "Number of alleles per locus")
> barplot(toto$Hexp - toto$Hobs, main = "Heterozygosity: expected-observed",
+      ylab = "Hexp - Hobs")
> barplot(toto$pop.eff, main = "Sample sizes per population", ylab = "Number of genotypes",
+      las = 3)
```



Is mean observed H significantly lower than mean expected H ?

```
> bartlett.test(list(toto$Hexp, toto$Hobs))
```

Bartlett test of homogeneity of variances

```
data: list(toto$Hexp, toto$Hobs)
```

```
Bartlett's K-squared = 0.047, df = 1, p-value = 0.8284
```

```
> t.test(toto$Hexp, toto$Hobs, pair = T, var.equal = TRUE, alter = "greater")
```

Paired t-test

```
data: toto$Hexp and toto$Hobs
```

```
t = 8.3294, df = 8, p-value = 1.631e-05
```

```
alternative hypothesis: true difference in means is greater than 0
```

```
95 percent confidence interval:
```

```
0.1134779      Inf
```

```
sample estimates:
```

```
mean of the differences
```

```
0.1460936
```

Yes, it is.

### 3.5 Measuring and testing population structure (a.k.a F statistics)

Population structure is traditionally measured and tested using F statistics, in particular  $F_{st}$ . *adeigenet* proposes different tools in this respect: general F statistics (`fstat`), a test of overall population structure (`gstat.randtest`), and pairwise  $F_{st}$  between all pairs of populations in a dataset (`pairwise.fst`). The first two are wrappers for functions implemented in the *hierfstat* package; pairwise  $F_{st}$  is implemented in *adeigenet*.

We illustrate their use using the dataset of microsatellite of cats from Nancy:

```
> library(hierfstat)
> data(nancycats)
> fstat(nancycats)
```

```

              pop      Ind
Total 0.08301274 0.1824701
pop    0.00000000 0.1084610
```

This table provides the three F statistics  $F_{st}$  (pop/total),  $F_{it}$  (Ind/total), and  $F_{is}$  (ind/pop). These are overall measures which take into account all genotypes and all loci.

Is the structure between populations significant? This question can be addressed using the G-statistic test (Goudet *et al.*, 1996); it is implemented for `genind` objects and produces a `randtest` object (package *ade4*).

```
> library(ade4)
> toto <- gstat.randtest(nancycats, nsim = 99)
> toto
```

```
Monte-Carlo test
Call: gstat.randtest(x = nancycats, nsim = 99)
```


```
Observation: 3416.974
```

```
Based on 99 replicates
Simulated p-value: 0.01
Alternative hypothesis: greater
```

```

      Std.Obs Expectation  Variance
29.69807  1776.07127  3052.87498
```

```
> plot(toto)
```



Yes, it is (the observed value is indicated on the right, while histograms correspond to the permuted values). Note that *hierfstat* allows for more elaborated tests, in particular when different levels of hierarchical clustering are available. Such tests are better done directly in *hierfstat*; for this, **genind** objects can be converted to the adequate format using **genind2hierfstat**. For instance:

```
> toto <- genind2hierfstat(nancycats)
> head(toto)
```

	pop	fca8	fca23	fca43	fca45	fca77	fca78	fca90	fca96	fca37
N215	1	NA	136146	139139	116120	156156	142148	199199	113113	208208
N216	1	NA	146146	139145	120126	156156	142148	185199	113113	208208
N217	1	135143	136146	141141	116116	152156	142142	197197	113113	210210
N218	1	133135	138138	139141	116126	150150	142148	199199	91105	208208
N219	1	133135	140146	141145	126126	152152	142148	193199	113113	208208
N220	1	135143	136146	145149	120126	150156	148148	193195	91113	208208

```
> varcomp.glob(toto$pop, toto[, -1])
```

```
$loc
      [,1]      [,2]      [,3]
fca8 0.08867161 0.116693199 0.6682028
fca23 0.05384247 0.077539920 0.6666667
fca43 0.05518935 0.066055996 0.6793249
fca45 0.05861271 -0.001026783 0.7083333
fca77 0.08810966 0.156863586 0.6329114
fca78 0.04869695 0.079006911 0.5654008
fca90 0.07540329 0.097194716 0.6497890
fca96 0.07538325 -0.005902071 0.7543860
fca37 0.04264094 0.116318729 0.4514768
```

```
$overall
      Pop      Ind      Error
```

```
0.5865502 0.7027442 5.7764917
```

```
$F
```

```

      Pop      Ind
Total 0.08301274 0.1824701
Pop    0.00000000 0.1084610
```

F statistics are provided in \$F; for instance, here,  $F_{st}$  is 0.083.

Lastly, pairwise  $Fst$  is frequently used as a measure of distance between populations. The function `pairwise.fst` computes Nei's estimator (Nei, 1973) of pairwise  $Fst$ , computed as:

$$Fst(A, B) = \frac{H_t - (n_A H_s(A) + n_B H_s(B)) / (n_A + n_B)}{H_t}$$

where A and B refer to the two populations of sample size  $n_A$  and  $n_B$  and respective expected heterozygosity  $H_s(A)$  and  $H_s(B)$ , and  $H_t$  is the expected heterozygosity in the whole dataset. For a given locus, expected heterozygosity is computed as  $1 - \sum p_i^2$ , where  $p_i$  is the frequency of the  $i$ th allele, and the  $\sum$  represents summation over all alleles. For multilocus data, the heterozygosity is simply averaged over all loci. These computations are achieved for all pairs of populations by the function `pairwise.fst`; we illustrate this on a subset of individuals of `nancycats` (computations for the whole dataset would take a few tens of seconds):

```
> matFst <- pairwise.fst(nancycats[1:50, treatOther = FALSE])
> matFst
```

```

      1      2      3
2 0.08018500
3 0.07140847 0.08200880
4 0.08163151 0.06512457 0.04131227
```

The resulting matrix is Euclidean when there are no missing values:

```
> is.euclid(matFst)
```

```
[1] TRUE
```

It can therefore be used in a Principal Coordinate Analysis (which requires Euclideanity), used to build trees, etc.

### 3.6 Testing for Hardy-Weinberg equilibrium

The Hardy-Weinberg equilibrium test is implemented for `genind` objects. The function to use is `HWE.test.genind`, and requires the package *genetics*. Here we first produce a matrix of p-values (`res="matrix"`) using parametric test. Monte Carlo procedure are more reliable but also more computer-intensive (use `permut=TRUE`).

```
> toto <- HWE.test.genind(nancycats, res = "matrix")
> dim(toto)
```

```
[1] 17 9
```

One test is performed per locus and population, *i.e.* 153 tests in this case. Thus, the first question is: which tests are highly significant?

```
> colnames(toto)
```

```
[1] "fca8" "fca23" "fca43" "fca45" "fca77" "fca78" "fca90" "fca96" "fca37"
```

```
> which(toto < 1e-04, TRUE)
```

```
      row col
P14   14   2
P02    2   7
P02    2   8
P05    5   9
```

Here, only 4 tests indicate departure from HW. Rows give populations, columns give markers. Now complete tests are returned, but the significant ones are already known.

```
> toto <- HWE.test.genind(nancycats, res = "full")
> toto$fca23$P06
```

```
Pearson's Chi-squared test
```

```
data:  tab
X-squared = 19.25, df = 10, p-value = 0.0372
```

```
> toto$fca90$P10
```

```
Pearson's Chi-squared test
```

```
data:  tab
X-squared = 19.25, df = 10, p-value = 0.0372
```

```
> toto$fca96$P10
```

```
Pearson's Chi-squared test
```

```
data: tab  
X-squared = 4.8889, df = 10, p-value = 0.8985
```

```
> toto$fca37$P13
```

```
Pearson's Chi-squared test
```

```
data: tab  
X-squared = 14.8281, df = 10, p-value = 0.1385
```

### 3.7 Performing a Principal Component Analysis on **genind** objects

The tables contained in **genind** objects can be submitted to a Principal Component Analysis (PCA) to seek a typology of individuals. Such analysis is straightforward using *ade4* to prepare data and *ade4* for the analysis *per se*. One has first to replace missing data. Putting each missing observation at the mean of the concerned allele frequency seems the best choice (NA will be stuck at the origin).

```
> data(microbov)  
> any(is.na(microbov$tab))
```

```
[1] TRUE
```

```
> sum(is.na(microbov$tab))
```

```
[1] 6325
```

There are 6325 missing data. Assuming that these are evenly distributed (for illustration purpose only!), we replace them using **na.replace**. As we intend to use a PCA, the appropriate replacement method is to put each NA at the mean of the corresponding allele (argument 'method' set to 'mean').


```
> obj <- na.replace(microbov, method = "mean")
```

```
Replaced 6325 missing values
```




Done. Now, the analysis can be performed. Data are centred but not scaled as 'units' are the same.

```
> pca1 <- dudi.pca(obj$tab, cent = TRUE, scale = FALSE, scannf = FALSE,
+   nf = 3)
> barplot(pca1$eig[1:50], main = "Eigenvalues")
```




Here we represent the genotypes and 95% inertia ellipses for populations.

```
> s.class(pca1$li, obj$pop, lab = obj$pop.names, sub = "PCA 1-2",
+   csub = 2)
> add.scatter.eig(pca1$eig[1:20], nf = 3, xax = 1, yax = 2, posi = "top")
```



This plane shows that the main structuring is between African and French breeds, the second structure reflecting genetic diversity among African breeds. The third axis reflects the diversity among French breeds: Overall, all breeds seem well differentiated.

```
> s.class(pca1$li, obj$pop, xax = 1, yax = 3, lab = obj$pop.names,
+         sub = "PCA 1-3", csub = 2)
> add.scatter.eig(pca1$eig[1:20], nf = 3, xax = 1, yax = 3, posi = "top")
```



### 3.8 Performing a Correspondance Analysis on **genpop** objects

Being contingency tables, the **@tab** in **genpop** objects can be submitted to a Correspondance Analysis (CA) to seek a typology of populations. The approach is very similar to the previous one for PCA. Missing data are first replaced during conversion from **genind**, but one could create a **genpop** with NAs and then use **na.replace** to get rid of missing observations.


```
> data(microbov)
> obj <- genind2genpop(microbov, missing = "chi2")
```

Converting data from a **genind** to a **genpop** object...

Replaced 0 missing values

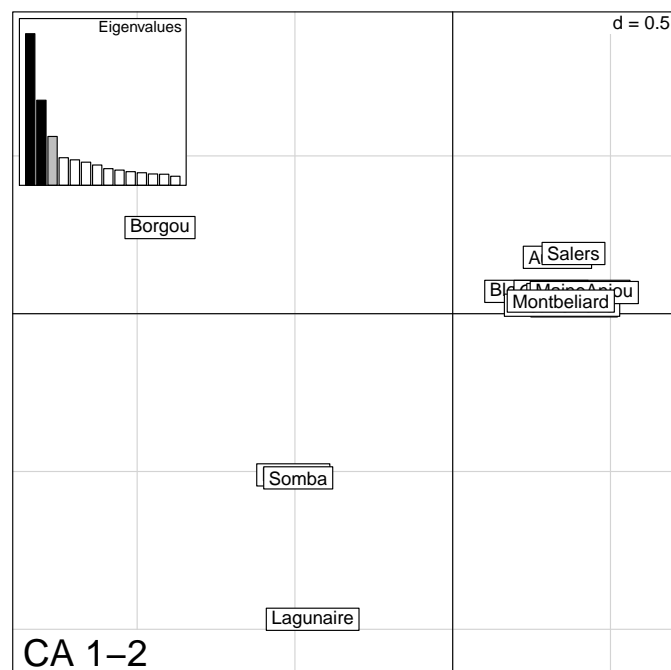
...done.

```
> ca1 <- dudi.coa(as.data.frame(obj$tab), scannf = FALSE, nf = 3)
> barplot(ca1$eig, main = "Eigenvalues")
```




Now we display the resulting typologies:

```
> s.label(ca1$li, lab = obj$pop.names, sub = "CA 1-2", csub = 2)
> add.scatter.eig(ca1$eig, nf = 3, xax = 1, yax = 2, posi = "top")
```



```
> s.label(ca1$li, xax = 1, yax = 3, lab = obj$pop.names, sub = "CA 1-3",
+         csub = 2)
> add.scatter.eig(ca1$eig, nf = 3, xax = 2, yax = 3, posi = "bottomright")
```



Once again, axes are to be interpreted separately in terms of continental differentiation, a among-breed diversities.


### 3.9 Analyzing a single locus

Here the emphasis is put on analyzing a single locus using different methods. Any marker can be isolated using the `seploc` instruction.

```
> data(nancycats)
> toto <- seploc(nancycats, truenames = TRUE, res.type = "matrix")
> X <- toto$fca90
```

`fca90.ind` is a matrix containing only genotypes for the marker `fca90`. It can be analyzed, for instance, using an inter-class PCA. This analysis provides a typology of individuals having maximal inter-colonies variance.

```
> library(ade4)
> pcaX <- dudi.pca(X, cent = T, scale = F, scannf = FALSE)
> pcabetX <- between(pcaX, nancycats$pop, scannf = FALSE)
> s.arrow(pcabetX$c1, xlim = c(-0.9, 0.9))
> s.class(pcabetX$l1, nancycats$pop, cell = 0, cstar = 0, add.p = T)
> sunflowerplot(X %*% as.matrix(pcabetX$c1), add = T)
> add.scatter.eig(pcabetX$eig, xax = 1, yax = 2, posi = "bottomright")
```



Here the differences between individuals are mainly expressed by three alleles: 199, 197 and 193. However, there is no clear structuration to be seen at an individual level. Is  $F_{st}$  significant taking only this marker into account? We perform the G-statistic test and eventually compute the corresponding F statistics. Note that we use the constructor **genind** to generate an object of this class from **X**:

```
> fca90.ind <- genind(X, pop = nancycats$pop)
> gstat.randtest(fca90.ind, nsim = 999)
```

```
Monte-Carlo test
Call: gstat.randtest(x = fca90.ind, nsim = 999)
```

```
Observation: 437.135
```

```
Based on 999 replicates
Simulated p-value: 0.001
Alternative hypothesis: greater
```

Std.Obs	Expectation	Variance
14.95085	188.47771	276.61161

```
> F <- varcomp(genind2hierfstat(fca90.ind))$F
> rownames(F) <- c("tot", "pop")
> colnames(F) <- c("pop", "ind")
> F
```

	pop	ind
tot	0.09168833	0.2098744
pop	0.00000000	0.1301162

In this case the information is best summarized by F statistics than by an ordination method. It is likely because all colonies are differentiated but none forming clusters of related colonies.

### 3.10 Testing for isolation by distance

Isolation by distance (IBD) is tested using Mantel test between a matrix of genetic distances and a matrix of geographic distances. It can be tested using individuals as well as populations. This example uses cat colonies. We use Edwards' distance *versus* Euclidean distances between colonies.

```
> data(nancycats)
> toto <- genind2genpop(nancycats, miss = "0")
```

```
Converting data from a genind to a genpop object...
```

```
Replaced 9 missing values
```

```
...done.
```

```
> Dgen <- dist.genpop(toto, method = 2)
> Dgeo <- dist(nancycats$other$xy)
> library(ade4)
> ibd <- mantel.randtest(Dgen, Dgeo)
> ibd
```

```
Monte-Carlo test
```

```
Call: mantel.randtest(m1 = Dgen, m2 = Dgeo)
```

```
Observation: 0.00492068
```


```
Based on 999 replicates
```

```
Simulated p-value: 0.487
```

```
Alternative hypothesis: greater
```

Std.Obs	Expectation	Variance
0.0003337896	0.0048859509	0.0108255438

```
> plot(ibd)
```



Isolation by distance is clearly not significant.

### 3.11 Using Monmonier's algorithm to define genetic boundaries

Monmonier's algorithm (Monmonier, 1973) was originally designed to find boundaries of maximum differences between contiguous polygons of a tessellation. As such, the method was basically used in geographical analysis. More recently, Manni *et al.* (2004) suggested that this algorithm could be employed to detect genetic boundaries among georeferenced genotypes (or populations). This algorithm is implemented using a more general approach than the initial one in `adegenet`.

Instead of using Voronoi tessellation as in original version, the functions `monmonier` and `optimize.monmonier` can handle various neighbouring graphs such as Delaunay triangulation, Gabriel's graph, Relative Neighbours graph, etc. These graphs defined spatial connectivity among 'points' (genotypes or populations), any couple of points being neighbours (if connected) or not. Another information is given by a set of markers which define genetic distances among these 'points'. The aim of Monmonier's algorithm is to find the path through the strongest genetic distances between neighbours. A more complete description of the principle of this algorithm will be found in the documentation of `monmonier`. Indeed, the very purpose of this tutorial is simply to show how it can be used on genetic data.

Let's take the example from the function's manpage and detail it. The dataset used is `sim2pop`.

```
> data(sim2pop)
> sim2pop

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:   genind
@call: old2new(object = sim2pop)

@tab: 130 x 241 matrix of genotypes

@ind.names: vector of 130 individual names
@loc.names: vector of 20 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 241 columns of @tab
@all.names: list of 20 components yielding allele names for each locus
@ploidy: 2
@type: codom


Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: xy
```

```
> summary(sim2pop$pop)
```

```
P01 P02  
100 30
```

```
> temp <- sim2pop$pop  
> levels(temp) <- c(17, 19)  
> temp <- as.numeric(as.character(temp))  
> plot(sim2pop$other$xy, pch = temp, cex = 1.5, xlab = "x", ylab = "y")  
> legend("topright", leg = c("Pop A", "Pop B"), pch = c(17, 19))
```



There are two sampled populations in this dataset, with unequal sample sizes (100 and 30). Twenty microsatellite-like loci are available for all genotypes (no missing data). So, what do **monmonier** ask for?

```
> args(monmonier)
```

```
function (xy, dist, cn, threshold = NULL, bd.length = NULL, nrun = 1,  
  skip.local.diff = rep(0, nrun), scanthres = is.null(threshold),  
  allowLoop = TRUE)  
NULL
```


The first argument (**xy**) is a matrix of geographic coordinates, already stored in **sim2pop**. Next argument is an object of class **dist**, which is basically a distance matrix cut in half. For now, we will use the classical Euclidean distance among alleles frequencies of genotypes. This is obtained by:



```
> D <- dist(sim2pop$tab)
```


The next argument (**cn**) is a connection network. As existing routines to build such networks are spread over several packages, the function **chooseCN** will help you choose one. This is an interactive function, so difficult to demonstrate here (see **?chooseCN**). Here we ask the function not to ask for a choice (**ask=FALSE**) and select the second type of graph which is the one of Gabriel (**type=2**).

```
> gab <- chooseCN(sim2pop$other$xy, ask = FALSE, type = 2)
```



The obtained network is automatically plotted by the function. It seems we are now ready to proceed to the algorithm.

```
> mon1 <- monmonier(sim2pop$other$xy, D, gab)
```



This plot shows all local differences sorted in decreasing order. The idea behind this is that a significant boundary would cause local differences to decrease abruptly after the boundary. This should be used to choose the *threshold* difference for the algorithm to stop. Here, no boundary is visible: we stop.

Why do the algorithm fail to find a boundary? Either because there is no genetic differentiation to be found, or because the signal differentiating both populations is too weak to overcome the random noise in genetic distances. What is the  $F_{st}$  between the two samples?

```
> library(hierfstat)
> temp <- genind2hierfstat(sim2pop)
> varcomp.glob(temp[, 1], temp[, -1])$F
```

```
      Pop      Ind
Total 0.03824374 -0.07541793
Pop    0.00000000 -0.11818137
```

This value is somewhat moderate ( $F_{st} = 0.038$ ). Is it significant?

```
> gtest <- gstat.randtest(sim2pop)
> gtest
```

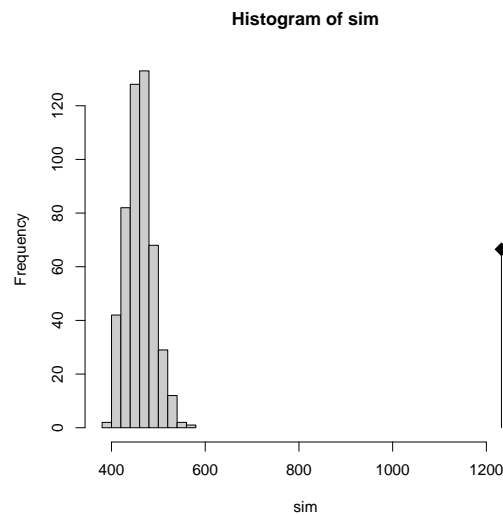
```
Monte-Carlo test
Call: gstat.randtest(x = sim2pop)
```

```
Observation: 1232.192
```

```
Based on 499 replicates
Simulated p-value: 0.002
Alternative hypothesis: greater
```


```
      Std.Obs Expectation  Variance
27.65661    459.81364    779.94176
```

```
> plot(gtest)
```



Yes, it is very significant. The two samples are indeed genetically differentiated. So, can Monmonier's algorithm find a boundary between the two populations? Yes, if we get rid of the random noise. This can be achieved using simple ordination method like Principal Coordinates Analysis.


```
> library(ade4)
> pco1 <- dudi.pco(D, scannf = FALSE, nf = 1)
> barplot(pco1$eig, main = "Eigenvalues")
```



We retain only the first eigenvalue. The corresponding coordinates are used to redefine the genetic distances among genotypes. The algorithm is then rerun.

```
> D <- dist(pco1$li)

> mon1 <- monmonier(sim2pop$other$xy, D, gab)
```



```
#####
# List of paths of maximum differences between neighbours #
#           Using a Monmonier based algorithm             #
#####

$call:monmonier(xy = sim2pop$other$xy, dist = D, cn = gab, scanthres = FALSE)

# Object content #
Class: monmonier
$run (number of successive runs): 1
$run1: run of the algorithm
$threshold (minimum difference between neighbours): 0.8154
$xy: spatial coordinates
$cn: connection network

# Runs content #
# Run 1
# First direction
Class: list
$path:
      x      y
Point_1 14.98299 93.81162

$values:
 2.281778
# Second direction
Class: list
$path:
      x      y
```

```
Point_1 14.98299 93.81162
Point_2 30.74508 87.57724
Point_3 33.66093 86.14115
...
```

```
$values:
2.281778 1.617905 1.953220 ...
```

This may take some time... but never more than five minutes on an 'ordinary' personal computer. The object `mon1` contains the whole information about the boundaries found. As several boundaries can be seeked at the same time (argument `nrun`), you have to specify about which run and which direction you want to get informations (values of differences or path coordinates). For instance:

```
> names(mon1)
```

```
[1] "run1"      "nrun"      "threshold" "xy"        "cn"        "call"
```

```
> names(mon1$run1)
```

```
[1] "dir1" "dir2"
```

```
> mon1$run1$dir1
```

```
$path
```

```
Point_1 14.98299 93.81162
```

```
$values
```

```
[1] 2.281778
```

It can also be useful to identify which points are crossed by the barrier; this can be done using `coords.monmonier`:

```
> coords.monmonier(mon1)
```

```
$run1
```

```
$run1$dir1
```

```
Point_1 14.98299 93.81162 11 125
```

```
$run1$dir2
```


```
Point_1 14.98299 93.81162 11 125
Point_2 30.74508 87.57724 44 128
Point_3 33.66093 86.14115 20 128
Point_4 35.28914 81.12578 68 128
Point_5 33.85756 74.45492 68 117
Point_6 38.07622 71.47532 68 122
Point_7 41.97494 70.02783 35 122
Point_8 43.45812 67.12026 69 122
```

Point_9	42.20206	59.59613	22	122
Point_10	42.48613	52.55145	22	124
Point_11	40.08702	48.61795	13	124
Point_12	39.20791	43.89978	13	127
Point_13	38.81236	40.34516	62	127
Point_14	37.32112	36.35265	62	130
Point_15	37.96426	30.82105	94	130
Point_16	32.79703	28.00517	16	130
Point_17	30.12832	28.60376	85	130
Point_18	20.92496	29.21211	63	119
Point_19	16.05811	22.72600	61	126
Point_20	11.72524	21.15519	89	126
Point_21	10.18696	16.61536	74	89

The returned dataframe contains, in this order, the  $x$  and  $y$  coordinates of the points of the barrier, and the identifiers of the two 'parent' points, that is, the points whose barycenter is the point of the barrier.


Finally, you can plot very simply the obtained boundary using the method `plot`:

```
> plot(mon1)
```



see arguments in `?plot.monmonier` to customize this representation. Last, we can compare the inferred boundary with the actual distribution of populations:

```
> plot(mon1, add.arrows = FALSE, bwd = 8)
> temp <- sim2pop$pop
> levels(temp) <- c(17, 19)
> temp <- as.numeric(as.character(temp))
> points(sim2pop$other$xy, pch = temp, cex = 1.3)
> legend("topright", leg = c("Pop A", "Pop B"), pch = c(17, 19))
```



Not too bad...

### 3.12 How to simulate hybridization?

The function `hybridize` allows to simulate hybridization between individuals from two distinct genetic pools, or more broadly between two `genind` objects. Here, we use the example from the manpage of the function, to go a little further. Please have a look at the documentation, especially at the different possible outputs (outputs for the software STRUCTURE is notably available).

```
> temp <- seppop(microbov)
> names(temp)
```

```
[1] "Borgou"      "Zebu"        "Lagunaire"   "NDama"
[5] "Somba"      "Aubrac"     "Bazadais"   "BlondeAquitaine"
[9] "BretPieNoire" "Charolais"  "Gascon"     "Limousin"
[13] "MaineAnjou" "Montbeliard" "Salers"
```

```
> salers <- temp$Salers
> zebu <- temp$Zebu
> zebler <- hybridize(salers, zebu, n = 40, pop = "zebler")
```


A first generation (F1) of hybrids 'zebler' is obtained. Is it possible to perform a backcross, say, with 'salers' population? Yes, here it is:

```
> F2 <- hybridize(salers, zebler, n = 40)
> F3 <- hybridize(salers, F2, n = 40)
> F4 <- hybridize(salers, F3, n = 40)
```

and so on... Are these hybrids still genetically distinct? Let's merge all hybrids in a single dataset and test for genetic differentiation:

```
> dat <- repool(zebler, F2, F3, F4)
> test <- gstat.randtest(dat)
> plot(test)
> temp <- genind2hierfstat(dat)
> varcomp.glob(temp[, 1], temp[, -1])$F
```

	Pop	Ind
Total	0.01384859	-0.03399172
Pop	0.00000000	-0.04851213



The  $F_{st}$  is not very strong (0.013) but still very significant: hybrids are still pretty well differentiated.

Finally, note that despite this example shows hybridization between diploid organisms, **hybridize** is not retrained to this case. In fact, organisms with any even level of ploidy can be used, in which case half of the genes is taken from each reference population. Ultimately, more complex mating schemes could be implemented... suggestion or (better) contributions are welcome!



### 3.13 Handling presence/absence data

Adegenet was primarily suited to handle codominant, multiallelic markers like microsatellites. However, dominant binary markers, like AFLP, can be used as well. In such a case, only presence/absence of alleles can be deduced accurately from the genotypes. This has several consequences, like the inability to compute allele frequencies. Hence, some functionalities in adegenet won't be available for dominant markers.

From version 1.2-3 of adegenet, the distinction between both types of markers is made by the slot 'type' of genind or genpop objects, which equals "codom" for codominant markers, and "PA" for presence/absence data. In the latter case, the 'tab' slot of a genind object no longer contains allele frequencies, but only presence/absence of alleles in a genotype. Similarly, the 'tab' slot of a genpop object no longer contains counts of alleles in the populations; instead, it contains the number of genotypes in each population possessing at least one copy of the concerned alleles. Moreover, in the case of presence/absence, the slots 'loc.nall', 'loc.fac', and 'all.names' become useless, and are thus all set to NULL.

Objects of type 'PA' are otherwise handled like usual (type 'codom') objects. Operations that are not available for PA type will issue an appropriate error message.

Here is an example using a toy dataset 'AFLP.txt' that can be downloaded from the adegenet website, section 'Documentation':

```
> dat <- read.table("http://adegenet.r-forge.r-project.org/files/AFLP.txt",
+   header = TRUE)
> dat
```

	loc1	loc2	loc3	loc4
indA	1	0	1	1
indB	0	1	1	1
indC	1	1	0	1
indD	0	NA	1	NA
indE	1	1	0	0
indF	1	0	1	1
indG	0	1	1	0

The function `df2genind` is used to obtain a genind object:

```
> obj <- genind(dat, ploidy = 1, type = "PA")
> obj
```

```
#####
### Genind object ###
#####
- genotypes of individuals -
S4 class: genind
```

```
@call: genind(tab = dat, ploidy = 1, type = "PA")
@tab: 7 x 4 matrix of genotypes
@ind.names: vector of 7 individual names
@loc.names: vector of 4 locus names
@loc.nall: NULL
@loc.fac: NULL
@all.names: NULL
@ploidy: 1
@type: PA

Optionnal contents:
@pop: - empty -
@pop.names: - empty -

@other: - empty -
```

```
> truenames(obj)
```

	loc1	loc2	loc3	loc4
indA	1	0	1	1
indB	0	1	1	1
indC	1	1	0	1
indD	0	NA	1	NA
indE	1	1	0	0
indF	1	0	1	1
indG	0	1	1	0

One can see that for instance, the summary of this object is more simple (no numbers of alleles per locus, no heterozygosity):

```
> pop(obj) <- rep(c("a", "b"), 4:3)
> summary(obj)
```

```
# Total number of genotypes: 7
# Population sample sizes:
a b
4 3
# Percentage of missing data:
[1] 7.142857
```

But we can still perform basic manipulation, like converting our object into a `genpop`:

```
> obj2 <- genind2genpop(obj)
```

```
Converting data from a genind to a genpop object...
...done.
```

```
> obj2
```

```
#####  
### Genpop object ###  
#####  
- Alleles counts for populations -  
  
S4 class:  genpop  
@call:  genind2genpop(x = obj)  
  
@tab:  2 x 4 matrix of alleles counts  
  
@pop.names: vector of  2 population names  
@loc.names: vector of  4 locus names  
@loc.nall: NULL  
@loc.fac: NULL  
@all.names: NULL  
@ploidy:  1  
@type:  PA  
  
@other: - empty -
```

```
> obj2@tab
```

```
  L1 L2 L3 L4  
1  2  2  3  3  
2  2  2  2  1
```

To continue with the toy example, we can proceed to a simple PCA. NAs are first replaced:

```
> objNoNa <- na.replace(obj, met = 0)
```


```
Replaced 2 missing values
```

```
> objNoNa@tab
```

```
  L1 L2 L3 L4  
1  1  0  1  1  
2  0  1  1  1  
3  1  1  0  1  
4  0  0  1  0  
5  1  1  0  0  
6  1  0  1  1  
7  0  1  1  0
```

Now the PCA is performed:

```
> library(ade4)  
> pca1 <- dudi.pca(objNoNa, scannf = FALSE, scale = FALSE)  
> scatter(pca1)
```



More generally, multivariate analyses from `ade4`, the sPCA (`spca`), the global and local tests (`global.rtest`, `local.rtest`), or the Monmonier's algorithm (`monmonier`) will work just fine with presence/absence data. However, it is clear that the usual Euclidean distance (used in PCA and sPCA), as well as many other distances, is not as accurate to measure genetic dissimilarity using presence/absence data as it is when using allele frequencies. The reason for this is that in presence/absence data, a part of the information is simply hidden. For instance, two individuals possessing the same allele will be considered at the same distance, whether they possess one or more copies of the allele. This might be especially problematic in organisms having a high degree of ploidy.

### 3.14 Assigning genotypes to clusters using Discriminant Analysis

The approach described below led to the development of a new methodological approach for studying the genetic diversity of biological populations, called the Discriminant Analysis of Principal Components (DAPC, Jombart et al. submitted). This method has been implemented by the functions `find.clusters` and `dapc` but is still considered under development. It will be documented along with this section pending the publication of the corresponding paper.

#### 3.14.1 Defining clusters

Bayesian clustering methods are not the only approaches for assigning genotypes to groups of genotypes. Discriminant analysis (DA; for a general presentation,

see Lachenbruch & Goldstein, 1979) is a multivariate method that has been used for the exact same purpose (Beharav & Nevo, 2003). It can be applied whenever pre-defined groups exist, to assign genotypes to and assess the robustness of these groups. New genotypes with unknown group can also be assigned to existing clusters. Although a few precautions have to be taken when applying DA (see Jombart *et al.* (2009) for a short overview), this is a useful and straightforward approach. It is here illustrated using cat colonies of Nancy, France (**nancycats** dataset).

```
> data(nancycats)
> nancycats

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:   genind
@call:   genind(tab = truenames(nancycats)$tab, pop = truenames(nancycats)$pop)

@tab:   237 x 108 matrix of genotypes

@ind.names: vector of 237 individual names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 108 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type: codom

Optionnal contents:
@pop: factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: xy

> unique(pop(nancycats))

[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
Levels: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
```

This dataset contains 237 genotypes of cats sampled over 17 colonies. A usual PCA on the allele frequencies of the populations would not show any structure, but colonies seem nonetheless mildly differentiated, as confirmed by Goudet's  $G$  test (and the  $F_{st}$  value):

```
> gstat.randtest(nancycats, n = 199)
```

```
Monte-Carlo test  
Call: gstat.randtest(x = nancycats, nsim = 199)
```

```
Observation: 3416.974
```

```
Based on 199 replicates  
Simulated p-value: 0.005  
Alternative hypothesis: greater
```

Std.Obs	Expectation	Variance
29.88014	1760.48172	3073.36041

```
> fstat(nancycats, fsonly = TRUE)
```

```
[1] 0.08301274
```

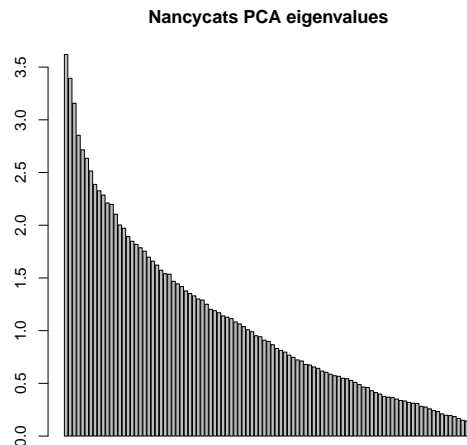
DA can be used to find the linear combinations of alleles that discriminate best the groups of genotypes (here, colonies). While a powerful method, DA is impaired by correlation between predictors, which arises for instance when linkage disequilibrium occurs between alleles. It is also impracticable when the number of alleles ( $p$ ) is greater than the number of genotypes ( $n$ ), and it generally requires  $n \gg p$  to yield reliable (numerically stable) results.

Thus, DA can seem often problematic when it comes to genetic data. One simple and efficient solution to all these issues is to transform alleles frequencies into a few independent (uncorrelated) components that summarise most of the genetic information, retaining only essential genetic features. This can be achieved by different multivariate methods; here, we shall use PCA. Genotype data are first transformed into scaled allele frequencies (using `scaleGen`):

```
> x <- scaleGen(nancycats, missing = "mean")
```

Then, we proceed to the PCA, retaining many principal components (PCs):

```
> x.pca <- dudi.pca(x, center = FALSE, scale = FALSE, scannf = FALSE,  
+   nf = 100)  
> barplot(x.pca$eig, main = "Nancycats PCA eigenvalues")
```



These eigenvalues indicate no structure, but this is no problem since here, we just use PCA as a mean of transforming genetic variables in an adequate way. PCs are stored in `x.pca$li`:

```
> head(x.pca$li[, 1:5])
```

	Axis1	Axis2	Axis3	Axis4	Axis5
N215	0.05847037	-0.4228593	0.4433385	-2.2204574	-1.94565586
N216	-0.30724734	-1.2376546	0.6085813	-0.8605435	-2.59138468
N217	0.12843698	1.6457391	1.4390692	-2.9874897	-0.09586744
N218	-0.93014696	-0.4694263	0.6595994	-2.0340478	-0.94786234
N219	-0.36584543	0.1139781	0.8042301	-0.9972465	-0.39733811
N220	-0.49235825	-1.2805118	0.2417573	-0.6060821	-1.63763566

```
> dim(x.pca$li)
```


```
[1] 237 99
```

Now, the question relates to how many PCs should be retained. This choice could be based on the success of assignment using DA (looking for the optimal value), or on a given fraction of the genetic diversity we would like to retain. We here use the latter, more simple and sufficient to illustrate the method. The following graph shows the cumulative amount of genetic information brought by each added PC:

```
> temp <- cumsum(x.pca$eig)/sum(x.pca$eig)
> plot(temp, xlab = "Added PC", ylab = "Fraction of the total genetic information")
> min(which(temp > 0.8))
```

[1] 52

```
> axis(1, at = 52, lab = 52)
> segments(52, 0, 52, temp[52], col = "red")
> segments(-5, 0.8, 52, 0.8, col = "red")
```



For instance, the first 52 PCs are sufficient to express 80% of the total genetic variability (see red segments). We choose to retain these 52 PCs and use them as new predictors in a DA. While there is a `discrimin` function in `ade4`, we use the function `lda` from the `MASS` package, which allows assigning (possibly new) genotypes to clusters.

```
> x.lda <- lda(x.pca$li[, 1:52], grouping = pop(nancycats))
> names(x.lda)
```

```
[1] "prior" "counts" "means" "scaling" "lev" "svd" "N"
[8] "call"
```

The object `x.lda` contains the results of the DA. For instance, coefficients of the linear combinations (*discriminant functions*) are stored in `x.lda$scaling`. For a further description of the content of these objects, see `?x.lda`. As far as assignment is concerned, the most interesting information is provided by `predict`:



```
> x.pred <- predict(x.lda)
> names(x.pred)
```

```
[1] "class"      "posterior" "x"
```

```
> x.pred$class
```

```

[1] 1 1 1 1 7 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 2
[26] 2 3 2 2 2 2 2 3 3 3 4 3 3 5 3 3 3 3 3 4 4 4
[51] 4 4 4 4 4 4 4 4 4 4 4 8 1 4 4 4 5 5 5 5 5 5
[76] 5 5 5 5 5 5 5 6 6 6 6 4 6 6 6 3 6 6 7 7 7 7
[101] 7 12 7 7 7 7 7 4 10 8 15 8 8 8 8 8 8 9 9 9 9
[126] 9 10 10 10 10 10 10 6 10 10 10 9 11 11 11 11 11 5 5 11 11 11 11
[151] 11 11 11 11 3 3 11 12 12 12 12 12 12 13 13 13 13 13 13 13 13 13
[176] 13 13 14 14 14 14 5 14 14 14 14 14 14 4 9 14 14 14 14 15 15 15 15
[201] 15 15 12 15 15 16 16 16 16 16 16 16 16 16 16 16 12 12 12 12 12 16
[226] 17 17 17 17 17 17 17 17 17 17 17 17 17 17
Levels: 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
```

```
> head(x.pred$posterior[, 1:5])
```

```

      1      2      3      4      5
N215 0.9789727 9.577822e-09 4.774675e-06 1.457425e-05 4.991132e-04
N216 0.8973141 1.927180e-08 2.356494e-06 2.098230e-04 5.954925e-05
N217 0.9999967 4.642204e-18 2.802290e-12 2.187224e-13 2.285345e-07
N218 0.8025053 6.453325e-12 2.245869e-12 1.445570e-12 7.946815e-10
N219 0.2111993 7.718517e-10 1.371071e-08 3.039278e-09 1.380308e-07
N220 0.9990817 9.608996e-10 1.596801e-07 1.499430e-07 4.625445e-04
```

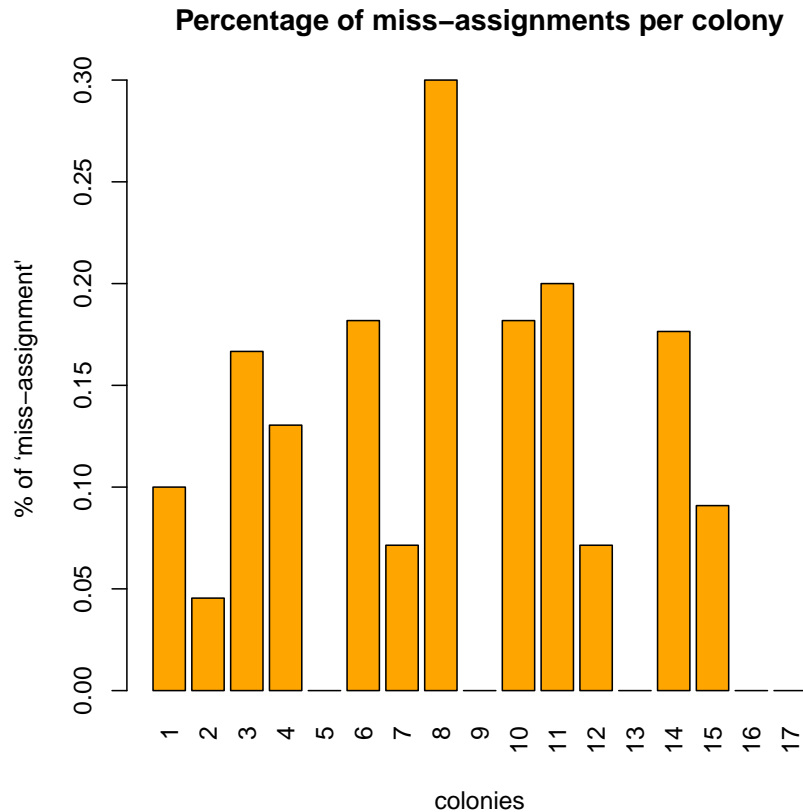
The `class` slot contains the cluster to which each genotype would be assigned with the highest probability, while `posterior` gives posterior probabilities of assignment of genotypes to clusters. The inferred groups can be compared easily to actual colonies:

```
> mean(x.pred$class == pop(nancycats))
```

```
[1] 0.8987342
```

In this case, each genotype would be assigned to the colony where it was actually found in 90% of cases. ‘Miss-assigned’ individuals could be hybrids or migrants, or simply reflect less clear-cut clusters. It is easy to check if some colonies have more of these:

```
> misAs <- tapply(x.pred$class != pop(nancycats), pop(nancycats),
+               mean)
> barplot(misAs, xlab = "colonies", ylab = "% of `miss-assignment'",
+         col = "orange", las = 3)
> title("Percentage of miss-assignments per colony")
```




For more details about genotypes, we can have a look at the **posterior** component, which gives probabilities of belonging to each cluster for each genotype:

```
> head(x.pred$posterior[, 1:10])
```

	1	2	3	4	5	6
N215	0.9789727	9.577822e-09	4.774675e-06	1.457425e-05	4.991132e-04	9.618366e-09
N216	0.8973141	1.927180e-08	2.356494e-06	2.098230e-04	5.954925e-05	2.773016e-09
N217	0.9999967	4.642204e-18	2.802290e-12	2.187224e-13	2.285345e-07	3.960112e-12
N218	0.8025053	6.453325e-12	2.245869e-12	1.445570e-12	7.946815e-10	1.027150e-13
N219	0.2111993	7.718517e-10	1.371071e-08	3.039278e-09	1.380308e-07	1.470112e-09
N220	0.9990817	9.608996e-10	1.596801e-07	1.499430e-07	4.625445e-04	2.407030e-10
	7	8	9	10		
N215	4.566426e-04	5.079117e-05	2.460440e-08	1.187296e-04		
N216	8.503242e-04	3.125840e-03	4.100707e-09	1.303221e-05		
N217	1.348738e-06	2.201991e-12	4.042410e-17	6.434196e-09		
N218	1.522997e-05	7.363011e-08	4.008925e-09	1.399387e-09		
N219	7.861739e-01	9.529260e-08	9.519421e-11	5.942611e-10		
N220	8.340154e-08	7.668008e-05	1.783684e-10	7.353064e-06		

This information is best perceived graphically (here, for the first 50 genotypes):

```
> table.paint(head(x.pred$posterior, 50), col.lab = paste("colony",
+ 1:17, sep = "."))
```



For instance, N215 (first row) is clearly assigned to colony 1, while it is unclear whether N158 (middle) belongs to colony 3 or 5. Such graphics is really good at summarising probabilities of assignment. In particular, it can be employed even when the number of clusters is relatively high, which would not be the case with classical graphs proposed in STRUCTURE.

### 3.14.2 Assigning new individuals

In certain cases, we may want to assign new genotypes to a pre-existing classification, as defined by a DA. This can be the case when new samples have been made available after a pilot study, or when doing cross-validation. We will simulate these cases by drawing 30 genotypes at random, and then trying to assign them to the defined clusters.

The following code only repeats the former analyses after withdrawing the 30 genotypes:

```
> id <- sample(1:237, 30)
> newSamp <- nancycats[id]
```

```
> newObj <- nancycats[-id]
> newx <- x[-id, ]
> newx.pca <- dudi.pca(newx, center = FALSE, scale = FALSE, scannf = FALSE,
+   nf = 100)
> newx.lda <- lda(newx.pca$li[, 1:52], grouping = pop(newObj))
```

The new object `x.lda` now contains a DA based on only 207 individuals. Our purpose is to assign the new genotypes (`newObj`) to existing clusters based on the discriminant functions defined in `x.lda`. It is a bit tricky, because we have to make sure new data are transformed like the old data, that is, with the same centring and scaling. Fortunately, centring and scaling values are stored as attributes in `x`, and can be provided to a new call to `scaleGen`:

```
> newSamp

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:   genind
@call:   .local(x = x, i = i, j = j, drop = drop)

@tab:   30 x 108 matrix of genotypes

@ind.names: vector of 30 individual names
@loc.names: vector of 9 locus names
@loc.nall: number of alleles per locus
@loc.fac: locus factor for the 108 columns of @tab
@all.names: list of 9 components yielding allele names for each locus
@ploidy: 2
@type:   codom

Optionnal contents:
@pop:   factor giving the population of each individual
@pop.names: factor giving the population of each individual

@other: a list containing: xy

> newSamp.x <- scaleGen(newSamp, center = attr(x, "scaled:center"),
+   scale = attr(x, "scaled:scale"), missing = "mean")

From there, the coordinates of these new genotypes onto the former PCA axes are
obtained easily using suprow:

> newSamp.pc <- suprow(newx.pca, newSamp.x)$lisup
> dim(newSamp.pc)

[1] 30 99

> head(newSamp.pc[, 1:5])
```


	Axis1	Axis2	Axis3	Axis4	Axis5
N84	0.5047608	0.04247984	0.9360310	-0.4568648	-0.5826832
N256	0.3689253	-0.83837998	0.2267852	0.6999959	1.2872675
N293	-0.2335015	1.11309930	0.4451031	-0.5755287	-0.6297878
N46	-0.3714081	1.64224571	-1.0602972	0.3324720	-0.4590934
N120	0.5303501	0.24115433	0.8207165	-0.7727013	0.8559355
N79	0.7176159	0.12113290	0.5928074	1.2419890	-0.3078119

The object `newSamp.pc` contains the supplementary principal components for the new 30 genotypes. To find the probabilities of belonging to any cluster for each new genotype, we simply use `predict`, and plot posterior probabilities as before, adding a green cross to indicate the actual colony:

```
> newSamp.pred <- predict(newx.lda, newSamp.pc[, 1:52])

> table.paint(newSamp.pred$posterior, col.lab = paste("colony",
+ 1:17, sep = "."))
> points(as.numeric(as.character(pop(newSamp))), 30:1, pch = "x",
+ col = "green", cex = 2)
> mean(as.character(newSamp.pred$class) == as.character(pop(newSamp)))
```

[1] 0.8



In this example, the new genotypes have been assigned to their actual group in 80% of cases. If our purpose was to cross-validate the classification of genotypes into groups, we would repeat this operation a large number of times, drawing a different random sample of genotypes each time.

## 4 Frequently Asked Questions

### 4.0.3 The function ... is not found. What's wrong?

You installed R, and adegenet, and all went ok. Yet, when trying to use some functions, like `read.genetix` for instance, you get an error message saying that the function is not found. The most likely explanation is that you do not have the most recent version of adegenet. This can be because you did not update your packages (see function `update.packages`). If your packages have been updated, and the problem persists, then you are likely using an outdated version of R, and though adegenet is up-to-date with respect to this R version, you are still using an outdated version of the package.

To know which version of adegenet you are using:

```
> packageDescription("adegenet", fields = "Version")
```

```
[1] "1.2-7"
```

And to know which version of R you are using:

```
> R.version.string
```

```
[1] "R version 2.11.1 (2010-05-31)"
```

## References

- BEHARAV, A. & NEVO, E. (2003). Predictive validity of discriminant analysis for genetic data. *Genetica* **119**, 259–267.
- CHARIF, D. & LOBRY, J. (2007). SeqinR 1.0-2: a contributed package to the R project for statistical computing devoted to biological sequences retrieval and analysis. In: *Structural approaches to sequence evolution: Molecules, networks, populations* (U. BASTOLLA, H. R., M. PORTO & VENDRUSCOLO, M., eds.), Biological and Medical Physics, Biomedical Engineering. New York: Springer Verlag, pp. 207–232. ISBN : 978-3-540-35305-8.

- CHESSEL, D., DUFOUR, A.-B. & THIOULOUSE, J. (2004). The ade4 package-I-one-table methods. *R News* **4**, 5–10.
- GOUDET, J. (2005). HIERFSTAT, a package for R to compute and test hierarchical F-statistics. *Molecular Ecology Notes* **5**, 184–186.
- GOUDET, J., RAYMOND, M., MEEÜS, T. & ROUSSET, F. (1996). Testing differentiation in diploid populations. *Genetics* **144**, 1933–1940.
- IHAKA, R. & GENTLEMAN, R. (1996). R: A language for data analysis and graphics. *Journal of Computational & Graphical Statistics* **5**, 299–314.
- JOMBART, T., PONTIER, D. & DUFOUR, A.-B. (2009). Genetic markers in the playground of multivariate analysis. *Heredity* **102**, 330–341.
- LACHENBRUCH, P. A. & GOLDSTEIN, M. (1979). Discriminant analysis. *Biometrics* **35**, 69–85.
- MANNI, F., GUÉRARD, E. & HEYER, E. (2004). Geographic patterns of (genetic, morphologic, linguistic) variation: how barriers can be detected by "Monmonier's algorithm". *Human Biology* **76**, 173–190.
- MONMONIER, M. (1973). Maximum-difference barriers: an alternative numerical regionalization method. *Geographical Analysis* **3**, 245–261.
- NEI, M. (1973). Analysis of gene diversity in subdivided populations. *Proc Natl Acad Sci U S A* **70**(12), 3321–3323.
- PARADIS, E. (2006). *Analysis of Phylogenetics and Evolution with R*. Springer-Verlag, Heidelberg.
- PARADIS, E., CLAUDE, J. & STRIMMER, K. (2004). APE: analyses of phylogenetics and evolution in R language. *Bioinformatics* **20**, 289–290.
- R DEVELOPMENT CORE TEAM (2009). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <http://www.R-project.org>. ISBN 3-900051-07-0.