

Practical course using the  software

Spatial genetics analyses using

Thibaut Jombart



Abstract

This practical course illustrates some methodological aspects of spatial genetics. In the following we shall test and describe spatial genetic structures, and compare the results obtained by different methods. A non-exhaustive list of packages used includes: *adeget* [8], *ade4* [4, 6, 5], *spdep*, *adehabitat* [2, 1], and *vegan* [10].

Contents

1	An overview of the data	3
2	Summarising the genetic diversity	4
3	Mapping and testing PCA results	9
4	Multivariate tests of spatial structure	14
5	spatial Principal Component Analysis	15

The chamois (*Rupicapra rupicapra*) is a conserved species in France. The Bauges mountains is a protected area in which the species has been recently studied. One of the most important questions for conservation purposes relates to whether individuals from this area form a single reproductive unit, or whether they are structured into sub-groups, and if so, what causes are likely to induce this structuring.

While field observations are very scarce and do not allow to answer this question, genetic data can be used to tackle the issue, as departure from panmixia should result in genetic structuring. The dataset *rupica* contains 335 georeferenced genotypes of Chamois from the Bauges mountains for 9 microsatellite markers, which we propose to analyse in this exercise.

1 An overview of the data

We first load the data:

```
> data(rupica)
> rupica

#####
### Genind object ###
#####
- genotypes of individuals -

S4 class:   genind
@call: NULL

@tab:  335 x 55 matrix of genotypes


@ind.names: vector of  335 individual names
@loc.names: vector of   9 locus names
@loc.nall:  number of alleles per locus
@loc.fac:  locus factor for the  55 columns of @tab
@all.names: list of   9 components yielding allele names for each locus
@ploidy:    2
@type:  codom

Optionnal contents:
@pop:  - empty -
@pop.names:  - empty -

@other: a list containing: xy  mnt  showBauges

rupica is a typical genind object, which is the class of objects storing
genotypes (as opposed to population data) in adegenet. rupica also contains
topographic information about the sampled area, which can be displayed by
calling rupica$other$showBauges. For instance, the spatial distribution of
the sampling can be displayed as follows:

> rupica$other$showBauges()
> points(rupica$other$xy, col = "red", pch = 20)
```




This spatial distribution is clearly not random, but seems arranged into loose clusters. However, superimposed samples can bias our visual assessment of the spatial clustering. Use a two-dimensional kernel density estimation (function `s.kde2d`) to overcome this possible issue.

Is geographical clustering strong enough to assign safely each individual to a group? Accordingly, shall we analyse these data at individual or group level?

2 Summarising the genetic diversity

As a prior clustering of genotypes is not known, we cannot employ usual F_{ST} -based approaches to detect genetic structuring. However, genetic structure could still result in a deficit of heterozygosity. Use the `summary` of `genind` objects to compare expected and observed heterozygosity:




The red line indicate identity between both quantities. What can we say about heterozygosity in this population? How can this be tested? The result below can be reproduced using a standard testing procedure:

Paired t-test

```
data: rupica.smry$Hexp and rupica.smry$Hobs
t = 0.9461, df = 8, p-value = 0.3718
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.01025068  0.02451318
sample estimates:
mean of the differences
      0.00713125
```

We can seek a global picture of the genetic diversity among genotypes using a Principal Component Analysis (PCA, [11, 7], `dudi.pca` in `ade4` package). The analysis is performed on a table of standardised alleles frequencies, obtained by `scaleGen` (use the binomial scaling option). Remember to disable the scaling option when performing the PCA. The function `dudi.pca` displays a barplot of eigenvalues and asks for a number of retained principal components:



The output produced by `dudi.pca` is a `dudi` object. A `dudi` object contains various information; in the case of PCA, principal axes (loadings), principal components (synthetic variable), and eigenvalues are respectively stored in `$c1`, `$li`, and `$eig` slots. Here is the content of the PCA:

```
Duality diagramm
class: pca dudi
$call: dudi.pca(df = rupica.X, center = FALSE, scale = FALSE, scannf = FALSE,
  nf = 2)

$nf: 2 axis-components saved
$rank: 45
eigen values: 1.561 1.34 1.168 1.097 1.071 ...
  vector length mode  content
1 $cw      55      numeric column weights
2 $lw     335      numeric row weights
3 $eig     45      numeric eigen values

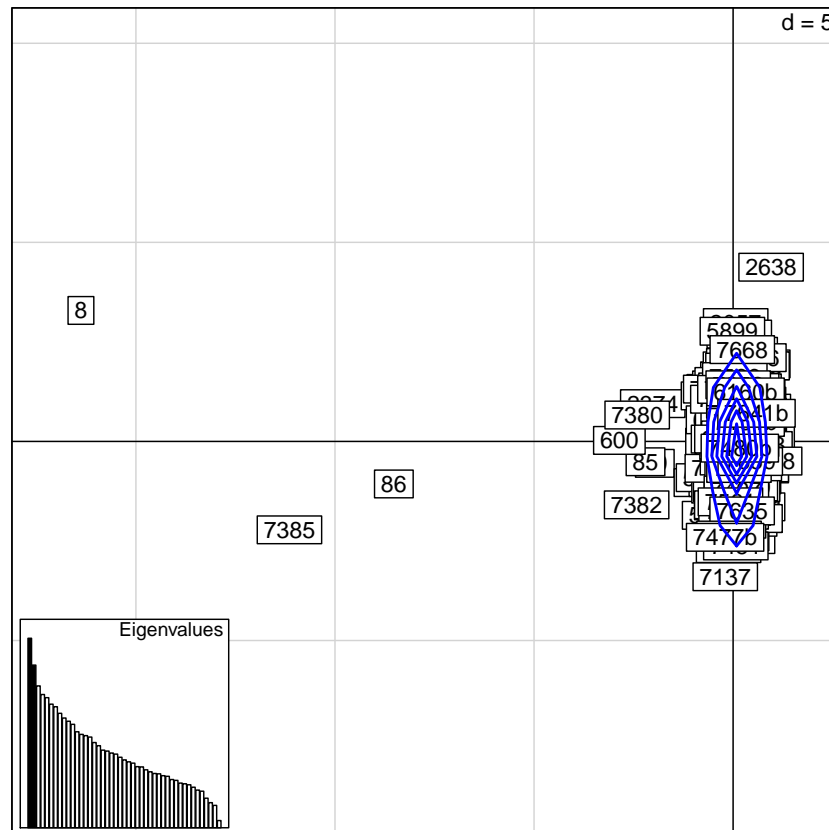
  data.frame nrow ncol content
1 $tab      335   55  modified array
2 $li       335    2   row coordinates
3 $l1       335    2   row normed scores
4 $co       55     2   column coordinates
5 $c1       55     2   column normed scores
other elements: cent norm
```

In general, eigenvalues represent the amount of genetic diversity — as measured by the multivariate method being used — represented by each principal component (PC). Verify that here, each eigenvalue is the variance of the corresponding PC.


An abrupt decrease in eigenvalues is likely to indicate the boundary between true patterns and non-interpretable structures. In this case, how

many PCs would you interpret?

Use `s.label` to display to two first components of the analysis. Then, use a kernel density (`s.kde2d`) for a better assessment of the distribution of the genotypes onto the principal axes:



What can we say about the genetic diversity among these genotypes as inferred by PCA? The function `loadingplot` allows to visualize the contribution of each allele, expressed as squared loadings, for a given principal component. Using this function, reproduce this figure:



What do we observe? We can get back to the genotypes for the concerned markers (e.g., Bm203) to check whether the highlighted genotypes are uncommon. `truenames` extracts the table of allele frequencies from a `genind` object (restoring original labels for markers, alleles, and individuals):

```
> X <- truenames(rupica)
> class(X)

[1] "matrix"

> dim(X)

[1] 335 55

> bm203.221 <- X[, "Bm203.221"]
> table(bm203.221)
```

```
bm203.221
      0 0.00597014925373134      0.5
    330      1      4
```

Only 4 genotypes possess one copy of the allele 221 of marker bm203 (the second result corresponds to a replaced missing data). Which individuals are they?


001 019 029 276
"8" "86" "600" "7385"

Conclusion?

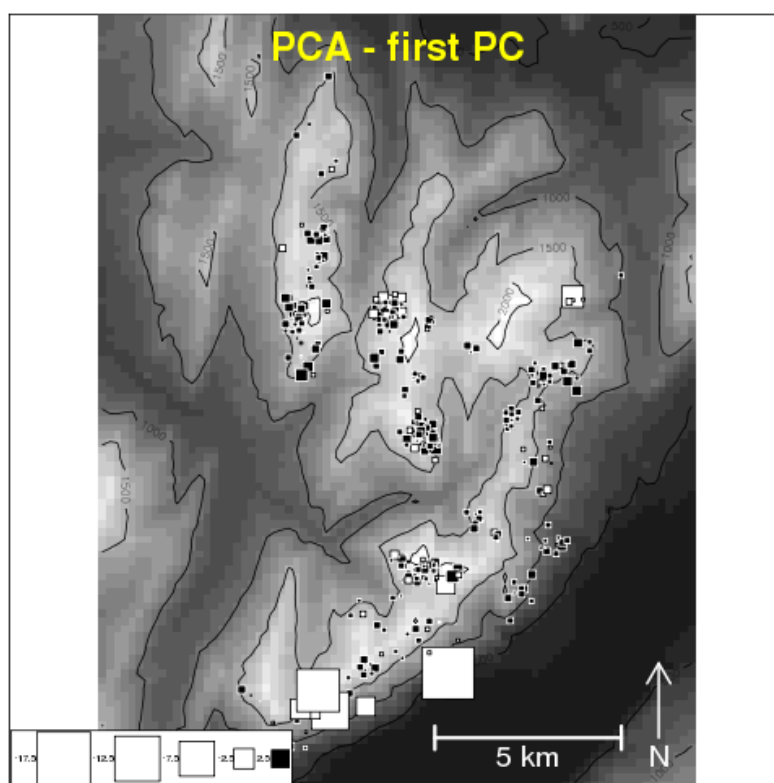
3 Mapping and testing PCA results

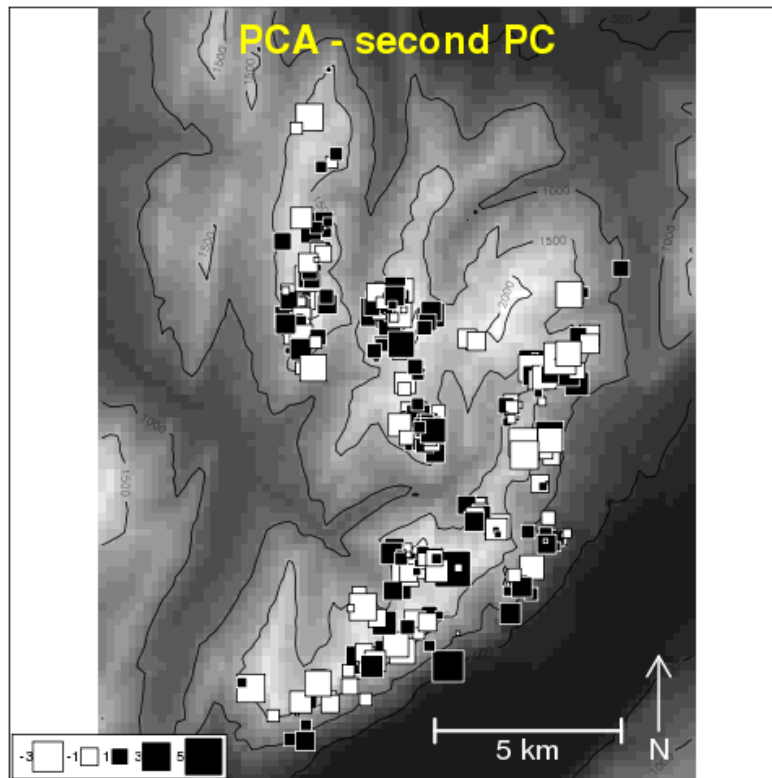
A frequent practice in spatial genetics is mapping the first principal components (PCs) onto the geographic space. The function `s.value` is well-suited to do so, using black and white squares of variable size for positive and negative values. To give a legend for this type of representation:

```
> s.value(cbind(1:11, rep(1, 11)), -5:5, cleg = 0)
> text(1:11, rep(1, 11), -5:5, col = "red", cex = 1.5)
```



Apply this graphical representation to the first two PCs of the PCA:





What can we say about spatial genetic structure as inferred by PCA? This visual assessment can be complemented by testing the spatial autocorrelation in the first PCs of PCA. This can be achieved using Moran's I test. Use the function `moran.mc` in the package `spdep` to perform these tests. You will need first to define the spatial connectivity between the sampled individuals. For these data, spatial connectivity is best defined as the overlap between home ranges of individuals. Home ranges will be modelled as disks with a radius of 1150m. Use `chooseCN` to create a connection network based on distance range ("neighbourhood by distance"). What threshold distance do you choose for individuals to be considered as neighbours? The connection network should resemble this:

```
Characteristics of weights list object:
Neighbour list object:
Number of regions: 335
Number of nonzero links: 18018
Percentage nonzero weights: 16.05525
Average number of links: 53.78507
```

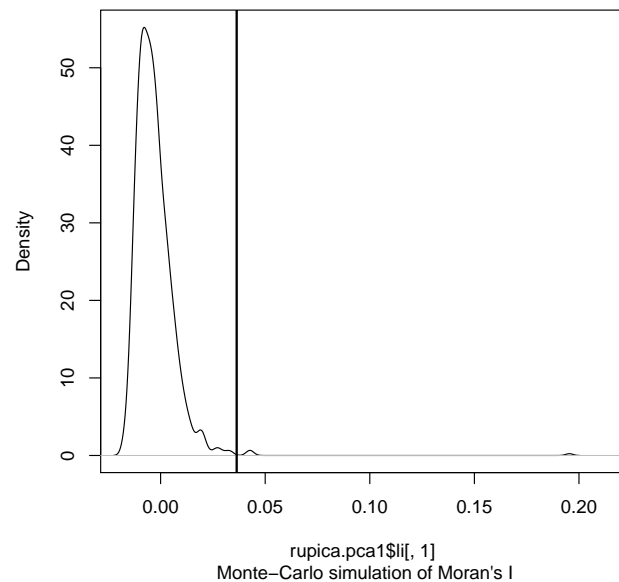
```
Weights style: W
Weights constants summary:
      n      nn  S0      S1      S2
W 335 112225 335 15.04311 1352.07
```

rupica.graph




Perform Moran's test for the first two PCs, and plot the results. The first test should be significant:

Density plot of permutation outcomes




Compare this result to the mapping of the first PC of PCA. What is wrong? When a test gives unexpected results, it is worth looking into the data in

more details. Moran's plot (`moran.plot`) plots the tested variable against its lagged vector. Use it on the first PC:




Actual positive autocorrelation corresponds to a positive correlation between a variable and its lag vector. Is it the case here? How can we explain that Moran's test was significant?

Repeat these analyses for the second PC. What are your conclusions?



4 Multivariate tests of spatial structure

So far, we have only tested the existence of spatial structures in the first two principal components of a PCA of the data. Therefore, these tests only describe one fragment of the data, and do not encompass the whole diversity in the data. As a complement, we can use Mantel test (`mantel.randtest`) to test spatial structures in the whole data, by assessing the correlation between genetic distances and geographic distances. Pairwise Euclidean distances are computed using `dist`. Perform Mantel test, using the scaled genetic data you used before in PCA, and the geographic coordinates.




What is your conclusion? Shall we be looking for spatial structures? If so, how can we explain that PCA did not reveal them? Does the Mantel correlogram (`mantel.correlog` in *vegan* package) bring any help solving the problem?

5 spatial Principal Component Analysis

The spatial Principal Component Analysis (sPCA, function `spca` [9]) has been especially developed to investigate hidden or non-obvious spatial genetic patterns. Like Moran's I test, sPCA first requires the spatial proximities between genotypes to be modeled. You will reuse the connection network defined previously using `chooseCN`, and pass it as the 'cn' argument of the function `spca`.

Read the documentation of `spca`, and apply the function to the dataset `rupica`. The function will display a barplot of eigenvalues:



This figure illustrates the fundamental difference between PCA and sPCA. Like `dudi.pca`, `spca` displays a barplot of eigenvalues, but unlike in PCA, eigenvalues of sPCA can also be negative. This is because the criterion optimized by the analysis can have positive and negative values, corresponding respectively to positive and negative autocorrelation. Positive spatial autocorrelation correspond to greater genetic similarity between geographically closer individuals. Conversely, negative spatial autocorrelation corresponds to greater dissimilarity between neighbours. The spatial autocorrelation of a variable is measured by Moran's I , and interpreted as follows:

- ★ $I_0 = -1/(n-1)$: no spatial autocorrelation (x is randomly distributed across space)
- ★ $I > I_0$: positive spatial autocorrelation
- ★ $I < I_0$: negative spatial autocorrelation

Principal components of PCA ensure that (ϕ referring to one PC) $var(\phi)$ is maximum. By contrast, sPCA provides PC which decompose the quantity $var(\phi)I(\phi)$. In other words, PCA focuses on variability only, while sPCA is a compromise between variability ($var(\phi)$) and spatial structure ($I(\phi)$).

In this case, only the principal components associated with the two first positive eigenvalues (in red) shall be retained. The printing of `spca` objects is more explicit than `dudi` objects, but named with the same conventions:

```
#####
# spatial Principal Component Analysis #
#####
class: spca
```



```
$call: spca(obj = rupica, cn = rupica.graph, scannf = FALSE, nfposi = 2,
  nfneg = 0)

$nfposi: 2 axis-components saved
$nfneg: 0 axis-components saved
Positive eigenvalues: 0.03018 0.01408 0.009211 0.006835 0.004529 ...
Negative eigenvalues: -0.008611 -0.006414 -0.004451 -0.003963 -0.003329 ...

  vector length mode  content
1 $eig    45      numeric eigenvalues


  data.frame nrow ncol content
1 $c1      55    2    principal axes: scaled vectors of alleles loadings
2 $li     335    2    principal components: coordinates of entities ('scores')
3 $ls     335    2    lag vector of principal components
4 $as      2    2    pca axes onto spca axes

$xy: matrix of spatial coordinates
$lw: a list of spatial weights (class 'listw')

other elements: NULL
```

Unlike usual multivariate analyses, eigenvalues of sPCA are composite: they measure both the genetic diversity (variance) and the spatial structure (spatial autocorrelation measured by Moran's I). This decomposition can also be used to choose which principal component to interpret. The function `screeplot` allows to display this information graphically:

Spatial and variance components of the eigenvalues



While λ_1 indicates with no doubt a structure, the second eigenvalue, λ_2 is

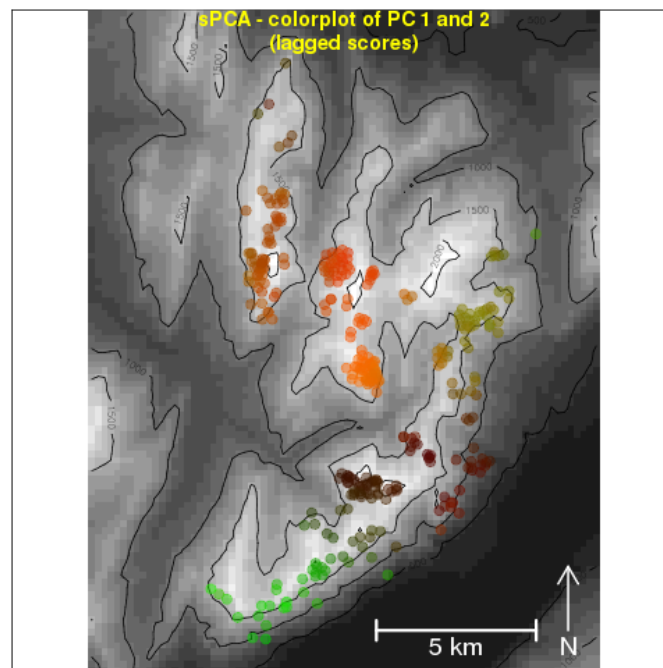
less clearly distinct from the successive values. Thus, we shall keep in mind this uncertainty when interpreting the second principal component of the analysis.

Try visualising the sPCA results as you did before with the PCA results. To clarify the possible spatial patterns, you can map the lagged PC ($\$ls$) instead of the PC ($\li), which are a 'denoised' version of the PCs.

First, map the first principal component of sPCA. How would you interpret this result? How does it compare to the first PC of PCA? What inference can we make about the way the landscape influences gene flow in this population of Chamois?

Do the same with the second PC of sPCA. Some field observations suggest that this pattern is not artefactual. How would you interpret this second structure?

To finish, you can try representing both structures at the same time using the color coding introduced by [3] (`?colorplot`). The final figure should resemble this (although colors may change from one computer to another):



References

- [1] C. Calenge. Exploring habitat selection by wildlife with adehabitat. *Journal of statistical software*, 22(6):1–19, 2007.
- [2] C. Calenge. The package "adehabitat" for the r software: a tool for the analysis of space and habitat use by animals. *Ecological Modelling*, 197:516–519, 2006.
- [3] L. L. Cavalli-Sforza, P. Menozzi, and A. Piazza. Demic expansions and human evolution. *Science*, 259:639–646, 1993.
- [4] D. Chessel, A.-B. Dufour, and J. Thioulouse. The ade4 package-I- one-table methods. *R News*, 4:5–10, 2004.
- [5] S. Dray and A.-B. Dufour. The ade4 package: implementing the duality diagram for ecologists. *Journal of Statistical Software*, 22(4):1–20, 2007.
- [6] S. Dray, A.-B. Dufour, and D. Chessel. The ade4 package - II: Two-table and K -table methods. *R News*, 7:47–54, 2007.
- [7] I. T. Joliffe. *Principal Component Analysis*. Springer Series in Statistics. Springer, second edition, 2004.
- [8] T. Jombart. adegenet: a R package for the multivariate analysis of genetic markers. *Bioinformatics*, 24:1403–1405, 2008.
- [9] T. Jombart, S. Devillard, A.-B. Dufour, and D. Pontier. Revealing cryptic spatial patterns in genetic variability by a new multivariate method. *Heredity*, 101:92–103, 2008.
- [10] Jari Oksanen, F. Guillaume Blanchet, Roeland Kindt, Pierre Legendre, R. B. O'Hara, Gavin L. Simpson, Peter Solymos, M. Henry H. Stevens, and Helene Wagner. *vegan: Community Ecology Package*, 2010. R package version 1.17-4.
- [11] K. Pearson. On lines and planes of closest fit to systems of points in space. *Philosophical Magazine*, 2:559–572, 1901.